

```
In [1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
```

```
In [74]: df=pd.read_csv(r"C:\Users\user\Downloads\csvs_per_year\csvs_per_year\madrid_2001.csv")
df
```

Out[74]:

| | date | BEN | CO | EBE | MXV | NMHC | NO_2 | NOx | OXY | O_3 | |
|--------|---------------------|-------|------|------|-------|------|-----------|------------|------|-----------|-------|
| 0 | 2001-08-01 01:00:00 | NaN | 0.37 | NaN | NaN | NaN | 58.400002 | 87.150002 | NaN | 34.529999 | 105.0 |
| 1 | 2001-08-01 01:00:00 | 1.50 | 0.34 | 1.49 | 4.10 | 0.07 | 56.250000 | 75.169998 | 2.11 | 42.160000 | 100.5 |
| 2 | 2001-08-01 01:00:00 | NaN | 0.28 | NaN | NaN | NaN | 50.660000 | 61.380001 | NaN | 46.310001 | 100.0 |
| 3 | 2001-08-01 01:00:00 | NaN | 0.47 | NaN | NaN | NaN | 69.790001 | 73.449997 | NaN | 40.650002 | 69.7 |
| 4 | 2001-08-01 01:00:00 | NaN | 0.39 | NaN | NaN | NaN | 22.830000 | 24.799999 | NaN | 66.309998 | 75.1 |
| ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... |
| 217867 | 2001-04-01 00:00:00 | 10.45 | 1.81 | NaN | NaN | NaN | 73.000000 | 264.399994 | NaN | 5.200000 | 47.8 |
| 217868 | 2001-04-01 00:00:00 | 5.20 | 0.69 | 4.56 | NaN | 0.13 | 71.080002 | 129.300003 | NaN | 13.460000 | 26.8 |
| 217869 | 2001-04-01 00:00:00 | 0.49 | 1.09 | NaN | 1.00 | 0.19 | 76.279999 | 128.399994 | 0.35 | 5.020000 | 40.7 |
| 217870 | 2001-04-01 00:00:00 | 5.62 | 1.01 | 5.04 | 11.38 | NaN | 80.019997 | 197.000000 | 2.58 | 5.840000 | 37.8 |
| 217871 | 2001-04-01 00:00:00 | 8.09 | 1.62 | 6.66 | 13.04 | 0.18 | 76.809998 | 206.300003 | 5.20 | 8.340000 | 35.3 |

217872 rows × 16 columns



```
In [3]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 217872 entries, 0 to 217871
Data columns (total 16 columns):
 #   Column      Non-Null Count  Dtype
---  -
 0   date        217872 non-null object
 1   BEN         70389 non-null float64
 2   CO          216341 non-null float64
 3   EBE         57752 non-null float64
 4   MXY         42753 non-null float64
 5   NMHC        85719 non-null float64
 6   NO_2        216331 non-null float64
 7   NOx         216318 non-null float64
 8   OXY         42856 non-null float64
 9   O_3         216514 non-null float64
10   PM10        207776 non-null float64
11   PXY         42845 non-null float64
12   SO_2        216403 non-null float64
13   TCH         85797 non-null float64
14   TOL         70196 non-null float64
15   station     217872 non-null int64
dtypes: float64(14), int64(1), object(1)
memory usage: 26.6+ MB
```

In [75]:

df=df.dropna()
df

Out[75]:

| | date | BEN | CO | EBE | MXY | NMHC | NO_2 | NOx | OXY | O_3 |
|--------|---------------------|-------|------|-------|-----------|------|------------|------------|-------|-----------|
| 1 | 2001-08-01 01:00:00 | 1.50 | 0.34 | 1.49 | 4.100000 | 0.07 | 56.250000 | 75.169998 | 2.11 | 42.160000 |
| 5 | 2001-08-01 01:00:00 | 2.11 | 0.63 | 2.48 | 5.940000 | 0.05 | 66.260002 | 118.099998 | 3.15 | 33.500000 |
| 21 | 2001-08-01 01:00:00 | 0.80 | 0.43 | 0.71 | 1.200000 | 0.10 | 27.190001 | 29.700001 | 0.76 | 56.990002 |
| 23 | 2001-08-01 01:00:00 | 1.29 | 0.34 | 1.41 | 3.090000 | 0.07 | 40.750000 | 51.570000 | 1.70 | 51.580002 |
| 25 | 2001-08-01 02:00:00 | 0.87 | 0.06 | 0.88 | 2.410000 | 0.01 | 29.709999 | 31.440001 | 1.20 | 56.520000 |
| ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... |
| 217829 | 2001-03-31 23:00:00 | 11.76 | 4.48 | 7.71 | 17.219999 | 0.89 | 103.900002 | 548.500000 | 7.62 | 9.680000 |
| 217847 | 2001-03-31 23:00:00 | 9.79 | 2.65 | 7.59 | 9.730000 | 0.46 | 91.320000 | 315.899994 | 3.75 | 6.660000 |
| 217849 | 2001-04-01 00:00:00 | 5.86 | 1.22 | 5.66 | 13.710000 | 0.25 | 64.370003 | 218.300003 | 6.46 | 7.480000 |
| 217853 | 2001-04-01 00:00:00 | 14.47 | 1.83 | 11.39 | 26.059999 | 0.33 | 84.230003 | 259.200012 | 11.39 | 5.440000 |
| 217871 | 2001-04-01 00:00:00 | 8.09 | 1.62 | 6.66 | 13.040000 | 0.18 | 76.809998 | 206.300003 | 5.20 | 8.340000 |

29669 rows × 16 columns

In [76]: `df.isnull().sum()`

```
Out[76]: date      0
BEN      0
CO       0
EBE      0
MXY      0
NMHC     0
NO_2     0
NOx      0
OXY      0
O_3      0
PM10     0
PXY      0
SO_2     0
TCH      0
TOL      0
station  0
dtype: int64
```

In [77]: `df.describe()`

```
Out[77]:
```

| | BEN | CO | EBE | MXY | NMHC | NO_2 | |
|--------------|--------------|--------------|--------------|--------------|--------------|--------------|-----|
| count | 29669.000000 | 29669.000000 | 29669.000000 | 29669.000000 | 29669.000000 | 29669.000000 | 296 |
| mean | 3.361895 | 1.005413 | 3.580229 | 8.113086 | 0.195222 | 67.652292 | 1 |
| std | 3.176669 | 0.863135 | 3.744496 | 7.909701 | 0.192585 | 34.003120 | 1 |
| min | 0.100000 | 0.000000 | 0.140000 | 0.210000 | 0.000000 | 1.180000 | |
| 25% | 1.280000 | 0.470000 | 1.390000 | 3.040000 | 0.080000 | 44.299999 | |
| 50% | 2.510000 | 0.760000 | 2.600000 | 5.830000 | 0.140000 | 64.449997 | 1 |
| 75% | 4.420000 | 1.270000 | 4.580000 | 10.640000 | 0.250000 | 86.540001 | 2 |
| max | 54.560001 | 11.890000 | 77.260002 | 150.600006 | 2.880000 | 292.700012 | 19 |

In [7]: `df.columns`

```
Out[7]: Index(['date', 'BEN', 'CO', 'EBE', 'MXY', 'NMHC', 'NO_2', 'NOx', 'OXY', 'O_3',
               'PM10', 'PXY', 'SO_2', 'TCH', 'TOL', 'station'],
              dtype='object')
```

```
In [8]: df.loc[df['NMHC']<1, 'NMHC']=0
df.loc[df['NMHC']>1, 'NMHC']=1
df['NMHC']=df['NMHC'].astype(int)
```

C:\ProgramData\Anaconda3\lib\site-packages\pandas\core\indexing.py:1720: SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame.

Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy (https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)

```
self._setitem_single_column(loc, value, pi)
```

C:\ProgramData\Anaconda3\lib\site-packages\pandas\core\indexing.py:1720: SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame.

Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy (https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)

```
self._setitem_single_column(loc, value, pi)
```

<ipython-input-8-c5145d14383f>:3: SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame.

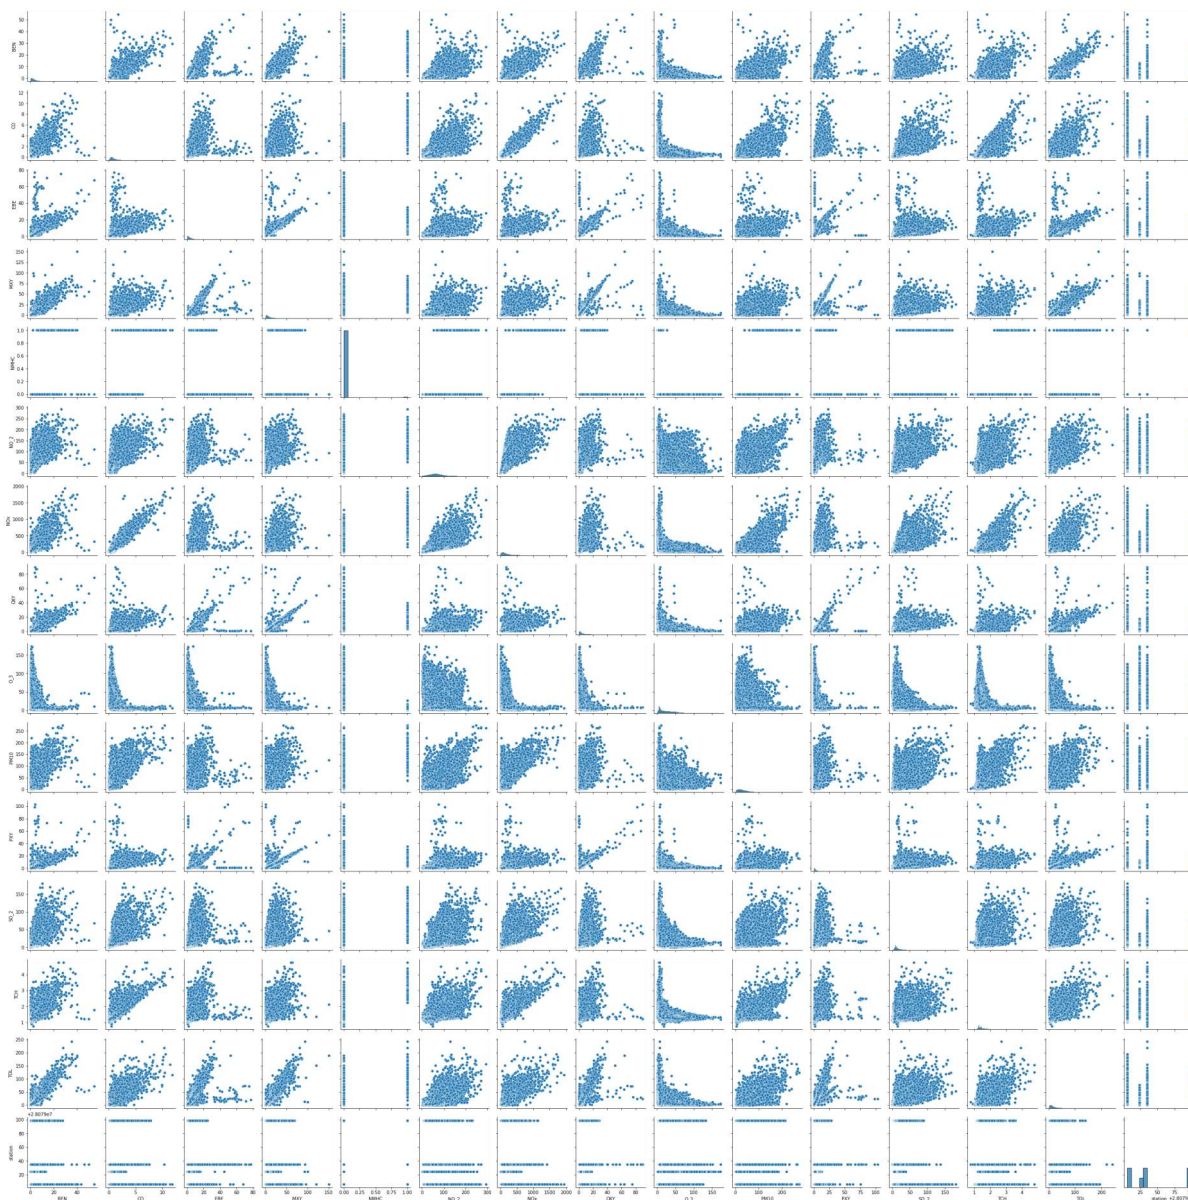
Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy (https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)

```
df['NMHC']=df['NMHC'].astype(int)
```

```
In [9]: sns.pairplot(df)
```

```
Out[9]: <seaborn.axisgrid.PairGrid at 0x1a131c8dfd0>
```




```
In [10]: sns.distplot(df['PM10'])
```

C:\ProgramData\Anaconda3\lib\site-packages\seaborn\distributions.py:2557: FutureWarning: `distplot` is a deprecated function and will be removed in a future version. Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).


```
warnings.warn(msg, FutureWarning)
```

```
Out[10]: <AxesSubplot:xlabel='PM10', ylabel='Density'>
```



```
In [11]: sns.heatmap(df.corr())
```

```
Out[11]: <AxesSubplot:>
```



LinearRegression

```
In [56]: x=df[['BEN', 'CO', 'EBE', 'MX', 'NO_2', 'NOx', 'OXY', 'O_3',  
            'station', 'PXY', 'SO_2', 'TCH', 'TOL']]  
y=df['PM10']  
from sklearn.model_selection import train_test_split  
x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.3)
```

```
In [57]: from sklearn.linear_model import LinearRegression  
lr=LinearRegression()  
lr.fit(x_train,y_train)
```

Out[57]: LinearRegression()

```
In [58]: print(lr.intercept_)
```

180436.05527078675

```
In [59]: coeff = pd.DataFrame(lr.coef_,x.columns,columns=['Co-efficient'])  
coeff
```


Out[59]:

| | Co-efficient |
|--|--------------|
|--|--------------|

| | |
|---------|-----------|
| BEN | -0.760819 |
| CO | -0.757628 |
| EBE | 0.562237 |
| MX | -0.593835 |
| NO_2 | 0.143722 |
| NOx | 0.050776 |
| OXY | -0.202122 |
| O_3 | 0.120388 |
| station | -0.006427 |
| PXY | 0.014732 |
| SO_2 | 0.115195 |
| TCH | 35.300122 |
| TOL | 0.549980 |


```
In [60]: prediction=lr.predict(x_test)
plt.scatter(y_test,prediction)
```

Out[60]: <matplotlib.collections.PathCollection at 0x1a14a3a7ee0>



```
In [61]: print(lr.score(x_test,y_test))
```

0.628018474497537

Ridge,Lasso

```
In [62]: from sklearn.linear_model import Ridge,Lasso
rr=Ridge(alpha=10)
rr.fit(x_train,y_train)
```

Out[62]: Ridge(alpha=10)

```
In [63]: rr.score(x_test,y_test)
```

Out[63]: 0.6279356506816863

```
In [64]: la=Lasso(alpha=10)
la.fit(x_train,y_train)
```

Out[64]: Lasso(alpha=10)

```
In [65]: la.score(x_test,y_test)
```

Out[65]: 0.5917301056999964

ElasticNet

```
In [66]: from sklearn.linear_model import ElasticNet
en=ElasticNet()
en.fit(x_train,y_train)
```

Out[66]: ElasticNet()

```
In [67]: print(en.coef_)
```

```
[-0.10220228  0.          0.32042828 -0.44626543  0.0994667   0.10047579
 -0.35859599  0.09424934  0.0240884   0.          0.06791359  0.53809407
  0.58111191]
```

```
In [68]: print(en.intercept_)
```

-676376.177026843

```
In [69]: print(en.predict(x_train))
```

```
[33.2024578  29.26592359 12.38848736 ... 54.60105985 36.04478949
 29.73010763]
```

```
In [70]: print(en.score(x_train,y_train))
```

0.6055408001396526

```
In [71]: from sklearn import metrics
print("Mean Absolytre Error:",metrics.mean_absolute_error(y_test,prediction))
```

Mean Absolytre Error: 12.577694478154767

```
In [72]: print("Mean Square Error:",metrics.mean_squared_error(y_test,prediction))
```

Mean Square Error: 316.8277572354703

```
In [73]: print("Root Mean Square Error:",np.sqrt(metrics.mean_absolute_error(y_test,prediction)))
```

Root Mean Square Error: 3.5465045436534783

LogisticRegression

```
In [30]: x=df[['BEN', 'CO', 'EBE', 'MXV', 'NO_2', 'NOx', 'OXY', 'O_3',  
            'PM10', 'PXY', 'SO_2', 'TCH','TOL', 'station']]  
y=df['NMHC']  
from sklearn.linear_model import LogisticRegression  
lgr=LogisticRegression()  
lgr.fit(x_train,y_train)
```

C:\ProgramData\Anaconda3\lib\site-packages\sklearn\linear_model_logistic.py:
763: ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max_iter) or scale the data as shown in:
<https://scikit-learn.org/stable/modules/preprocessing.html> (<https://scikit-learn.org/stable/modules/preprocessing.html>)
Please also refer to the documentation for alternative solver options:
https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression (https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression)
n_iter_i = _check_optimize_result(

Out[30]: LogisticRegression()

```
In [31]: lgr.predict(x_test)
```

Out[31]: array([28079099, 28079035, 28079035, ..., 28079035, 28079024, 28079006],
dtype=int64)

```
In [32]: lgr.score(x_test,y_test)
```

Out[32]: 0.6304909560723514

```
In [33]: from sklearn.preprocessing import StandardScaler  
fs=StandardScaler().fit_transform(x)  
logr=LogisticRegression()  
logr.fit(fs,y)
```

Out[33]: LogisticRegression()

```
In [34]: o=[[1,2,3,4,5,6,7,8,9,10,11,12,13,14]]  
prediction=logr.predict(o)  
print(prediction)
```

[0]

```
In [35]: logr.classes_
```

Out[35]: array([0, 1])

```
In [36]: logr.predict_proba(o)[0][0]
```

Out[36]: 0.9999926193823634

```
In [37]: logr.predict_proba(o)[0][1]
```

```
Out[37]: 7.3806176366125295e-06
```

RandomForest

```
In [48]: from sklearn.ensemble import RandomForestClassifier
rfc=RandomForestClassifier()
rfc.fit(x_train,y_train)
```

```
Out[48]: RandomForestClassifier()
```

```
In [49]: parameters={'max_depth':[1,2,3,4,5],
                    'min_samples_leaf':[5,10,15,20,25],
                    'n_estimators':[10,20,30,40,50]}
```

```
In [50]: from sklearn.model_selection import GridSearchCV
grid_search=GridSearchCV(estimator=rfc,param_grid=parameters,cv=2,scoring="accuracy")
grid_search.fit(x_train,y_train)
```

```
Out[50]: GridSearchCV(cv=2, estimator=RandomForestClassifier(),
                    param_grid={'max_depth': [1, 2, 3, 4, 5],
                                'min_samples_leaf': [5, 10, 15, 20, 25],
                                'n_estimators': [10, 20, 30, 40, 50]},
                    scoring='accuracy')
```

```
In [51]: grid_search.best_score_
```

```
Out[51]: 0.7233724961479199
```

```
In [52]: rfc_best=grid_search.best_estimator_
```

```
In [55]: from sklearn.tree import plot_tree  
plt.figure(figsize=(80,40))  
plot_tree(rfc_best.estimators_[5],class_names=['Yes','No','Yes','No'],filled=True)
```

```

Out[55]: [Text(2352.9, 1993.2, 'X[10] <= 19.67\ngini = 0.735\nsamples = 13118\nvalue =
[5899, 3001, 5967, 5901]\nclass = Yes'),
Text(1190.4, 1630.8000000000002, 'X[3] <= 1.355\ngini = 0.73\nsamples = 8929
\nvalue = [2305, 2921, 4857, 4074]\nclass = Yes'),
Text(595.2, 1268.4, 'X[5] <= 21.83\ngini = 0.464\nsamples = 1292\nvalue = [1
0, 1422, 493, 140]\nclass = No'),
Text(297.6, 906.0, 'X[1] <= 0.115\ngini = 0.177\nsamples = 629\nvalue = [2,
924, 68, 28]\nclass = No'),
Text(148.8, 543.5999999999999, 'X[6] <= 0.465\ngini = 0.307\nsamples = 42\nv
alue = [0, 14, 60, 0]\nclass = Yes'),
Text(74.4, 181.19999999999982, 'gini = 0.145\nsamples = 22\nvalue = [0, 3, 3
5, 0]\nclass = Yes'),
Text(223.20000000000002, 181.19999999999982, 'gini = 0.424\nsamples = 20\nva
lue = [0, 11, 25, 0]\nclass = Yes'),
Text(446.40000000000003, 543.5999999999999, 'X[11] <= 1.235\ngini = 0.078\ns
amples = 587\nvalue = [2, 910, 8, 28]\nclass = No'),
Text(372.0, 181.19999999999982, 'gini = 0.657\nsamples = 24\nvalue = [2, 7,
8, 16]\nclass = No'),
Text(520.80000000000001, 181.19999999999982, 'gini = 0.026\nsamples = 563\nva
lue = [0, 903, 0, 12]\nclass = No'),
Text(892.80000000000001, 906.0, 'X[11] <= 1.265\ngini = 0.594\nsamples = 663
\nvalue = [8, 498, 425, 112]\nclass = No'),
Text(744.0, 543.5999999999999, 'X[10] <= 7.815\ngini = 0.415\nsamples = 245
\nvalue = [8, 26, 286, 65]\nclass = Yes'),
Text(669.6, 181.19999999999982, 'gini = 0.219\nsamples = 166\nvalue = [5, 2
0, 234, 7]\nclass = Yes'),
Text(818.40000000000001, 181.19999999999982, 'gini = 0.568\nsamples = 79\nval
ue = [3, 6, 52, 58]\nclass = No'),
Text(1041.60000000000001, 543.5999999999999, 'X[10] <= 8.985\ngini = 0.436\ns
amples = 418\nvalue = [0, 472, 139, 47]\nclass = No'),
Text(967.2, 181.19999999999982, 'gini = 0.555\nsamples = 69\nvalue = [0, 36,
63, 11]\nclass = Yes'),
Text(1116.0, 181.19999999999982, 'gini = 0.343\nsamples = 349\nvalue = [0, 4
36, 76, 36]\nclass = No'),
Text(1785.60000000000001, 1268.4, 'X[7] <= 5.375\ngini = 0.713\nsamples = 763
7\nvalue = [2295, 1499, 4364, 3934]\nclass = Yes'),
Text(1488.0, 906.0, 'X[0] <= 2.39\ngini = 0.288\nsamples = 462\nvalue = [68,
615, 11, 41]\nclass = No'),
Text(1339.2, 543.5999999999999, 'X[6] <= 2.655\ngini = 0.054\nsamples = 155
\nvalue = [3, 248, 3, 1]\nclass = No'),
Text(1264.80000000000002, 181.19999999999982, 'gini = 0.027\nsamples = 135\nv
alue = [2, 218, 0, 1]\nclass = No'),
Text(1413.60000000000001, 181.19999999999982, 'gini = 0.213\nsamples = 20\nva
lue = [1, 30, 3, 0]\nclass = No'),
Text(1636.80000000000002, 543.5999999999999, 'X[10] <= 14.2\ngini = 0.39\nsam
ples = 307\nvalue = [65, 367, 8, 40]\nclass = No'),
Text(1562.4, 181.19999999999982, 'gini = 0.2\nsamples = 213\nvalue = [29, 29
1, 1, 6]\nclass = No'),
Text(1711.2, 181.19999999999982, 'gini = 0.646\nsamples = 94\nvalue = [36, 7
6, 7, 34]\nclass = No'),
Text(2083.20000000000003, 906.0, 'X[3] <= 8.715\ngini = 0.691\nsamples = 7175
\nvalue = [2227, 884, 4353, 3893]\nclass = Yes'),
Text(1934.4, 543.5999999999999, 'X[2] <= 0.895\ngini = 0.677\nsamples = 5738
\nvalue = [1291, 836, 3446, 3552]\nclass = No'),
Text(1860.00000000000002, 181.19999999999982, 'gini = 0.46\nsamples = 589\nva
lue = [21, 103, 647, 142]\nclass = Yes'),
Text(2008.80000000000002, 181.19999999999982, 'gini = 0.68\nsamples = 5149\nv

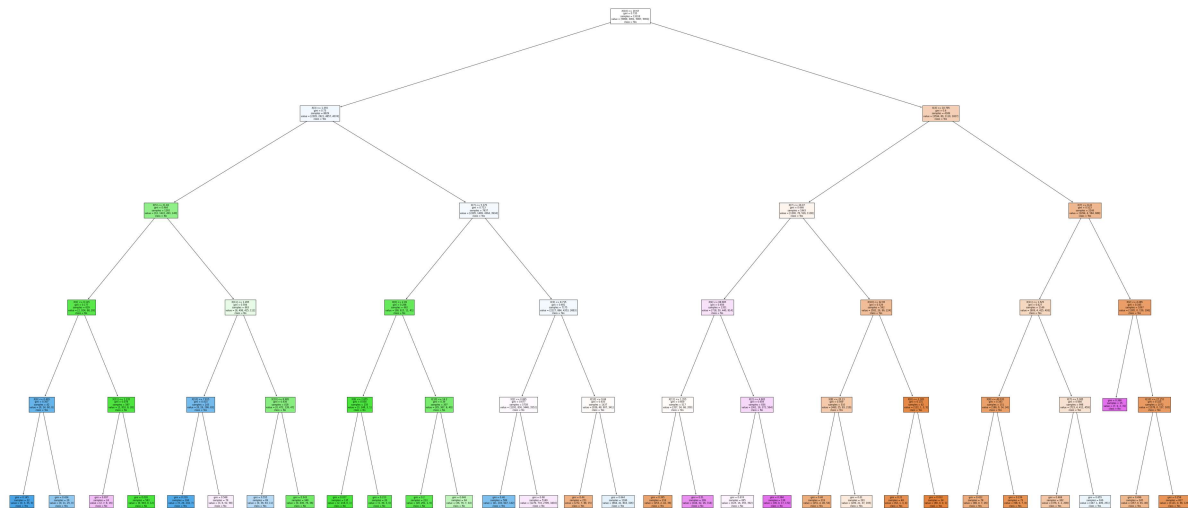
```

```

value = [1270, 733, 2799, 3410]\n\nclass = No'),
  Text(2232.0, 543.5999999999999, 'X[10] <= 9.66\ngini = 0.635\nsamples = 1437\n\nvalue = [936, 48, 907, 341]\n\nclass = Yes'),
  Text(2157.6000000000004, 181.19999999999982, 'gini = 0.44\nsamples = 251\n\nvalue = [272, 7, 89, 15]\n\nclass = Yes'),
  Text(2306.4, 181.19999999999982, 'gini = 0.644\nsamples = 1186\n\nvalue = [664, 41, 818, 326]\n\nclass = Yes'),
  Text(3515.4, 1630.8000000000002, 'X[3] <= 10.795\ngini = 0.6\nsamples = 4189\n\nvalue = [3594, 80, 1110, 1827]\n\nclass = Yes'),
  Text(2976.0, 1268.4, 'X[7] <= 26.07\ngini = 0.646\nsamples = 1943\n\nvalue = [1300, 76, 526, 1138]\n\nclass = Yes'),
  Text(2678.4, 906.0, 'X[4] <= 68.665\ngini = 0.656\nsamples = 1351\n\nvalue = [718, 50, 440, 914]\n\nclass = No'),
  Text(2529.6000000000004, 543.5999999999999, 'X[11] <= 1.355\ngini = 0.609\nsamples = 517\n\nvalue = [357, 34, 68, 350]\n\nclass = Yes'),
  Text(2455.2000000000003, 181.19999999999982, 'gini = 0.395\nsamples = 218\n\nvalue = [253, 2, 42, 36]\n\nclass = Yes'),
  Text(2604.0, 181.19999999999982, 'gini = 0.51\nsamples = 299\n\nvalue = [104, 32, 26, 314]\n\nclass = No'),
  Text(2827.2000000000003, 543.5999999999999, 'X[2] <= 4.665\ngini = 0.659\nsamples = 834\n\nvalue = [361, 16, 372, 564]\n\nclass = No'),
  Text(2752.8, 181.19999999999982, 'gini = 0.674\nsamples = 695\n\nvalue = [325, 16, 355, 392]\n\nclass = No'),
  Text(2901.6000000000004, 181.19999999999982, 'gini = 0.384\nsamples = 139\n\nvalue = [36, 0, 17, 172]\n\nclass = No'),
  Text(3273.6000000000004, 906.0, 'X[10] <= 32.65\ngini = 0.529\nsamples = 592\n\nvalue = [582, 26, 86, 224]\n\nclass = Yes'),
  Text(3124.8, 543.5999999999999, 'X[8] <= 20.13\ngini = 0.569\nsamples = 510\n\nvalue = [460, 25, 83, 219]\n\nclass = Yes'),
  Text(3050.4, 181.19999999999982, 'gini = 0.46\nsamples = 219\n\nvalue = [251, 4, 46, 54]\n\nclass = Yes'),
  Text(3199.2000000000003, 181.19999999999982, 'gini = 0.61\nsamples = 291\n\nvalue = [209, 21, 37, 165]\n\nclass = Yes'),
  Text(3422.4, 543.5999999999999, 'X[6] <= 3.185\ngini = 0.131\nsamples = 82\n\nvalue = [122, 1, 3, 5]\n\nclass = Yes'),
  Text(3348.0000000000005, 181.19999999999982, 'gini = 0.21\nsamples = 48\n\nvalue = [62, 1, 3, 4]\n\nclass = Yes'),
  Text(3496.8, 181.19999999999982, 'gini = 0.032\nsamples = 34\n\nvalue = [60, 0, 0, 1]\n\nclass = Yes'),
  Text(4054.8, 1268.4, 'X[7] <= 8.41\ngini = 0.523\nsamples = 2246\n\nvalue = [294, 4, 584, 689]\n\nclass = Yes'),
  Text(3868.8, 906.0, 'X[11] <= 1.525\ngini = 0.627\nsamples = 1149\n\nvalue = [909, 4, 425, 493]\n\nclass = Yes'),
  Text(3720.0000000000005, 543.5999999999999, 'X[8] <= 40.035\ngini = 0.343\nsamples = 151\n\nvalue = [186, 0, 14, 34]\n\nclass = Yes'),
  Text(3645.6000000000004, 181.19999999999982, 'gini = 0.422\nsamples = 78\n\nvalue = [88, 0, 7, 26]\n\nclass = Yes'),
  Text(3794.4, 181.19999999999982, 'gini = 0.239\nsamples = 73\n\nvalue = [98, 0, 7, 8]\n\nclass = Yes'),
  Text(4017.6000000000004, 543.5999999999999, 'X[7] <= 5.385\ngini = 0.646\nsamples = 998\n\nvalue = [723, 4, 411, 459]\n\nclass = Yes'),
  Text(3943.2000000000003, 181.19999999999982, 'gini = 0.468\nsamples = 362\n\nvalue = [376, 3, 2, 208]\n\nclass = Yes'),
  Text(4092.0000000000005, 181.19999999999982, 'gini = 0.655\nsamples = 636\n\nvalue = [347, 1, 409, 251]\n\nclass = Yes'),
  Text(4240.8, 906.0, 'X[2] <= 4.285\ngini = 0.345\nsamples = 1097\n\nvalue = [1385, 0, 159, 196]\n\nclass = Yes'),

```

```
Text(4166.400000000001, 543.5999999999999, 'gini = 0.366\nsamples = 25\nvalue = [7, 0, 2, 31]\nclass = No'),
Text(4315.200000000001, 543.5999999999999, 'X[10] <= 27.155\ngini = 0.325\nsamples = 1072\nvalue = [1378, 0, 157, 165]\nclass = Yes'),
Text(4240.8, 181.1999999999982, 'gini = 0.496\nsamples = 245\nvalue = [257, 0, 91, 40]\nclass = Yes'),
Text(4389.6, 181.1999999999982, 'gini = 0.258\nsamples = 827\nvalue = [112 1, 0, 66, 125]\nclass = Yes')]
```



```
In [ ]: Best model:LogisticRegression
```