


```
In [1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.linear_model import Ridge,Lasso
from sklearn.linear_model import ElasticNet
from sklearn import metrics
from sklearn.linear_model import LogisticRegression
from sklearn.preprocessing import StandardScaler
from sklearn.ensemble import RandomForestClassifier
from sklearn.model_selection import GridSearchCV
from sklearn.tree import plot_tree
```

```
In [2]: df=pd.read_csv(r"C:\Users\user\Downloads\csvs_per_year\csvs_per_year\madrid_2014.csv")
df
```

Out[2]:

	date	BEN	CO	EBE	NMHC	NO	NO_2	O_3	PM10	PM25	SO_2	TCH	TOL	
0	2014-06-01 01:00:00	NaN	0.2	NaN	NaN	3.0	10.0	NaN	NaN	NaN	3.0	NaN	NaN	28
1	2014-06-01 01:00:00	0.2	0.2	0.1	0.11	3.0	17.0	68.0	10.0	5.0	5.0	1.36	1.3	28
2	2014-06-01 01:00:00	0.3	NaN	0.1	NaN	2.0	6.0	NaN	NaN	NaN	NaN	NaN	1.1	28
3	2014-06-01 01:00:00	NaN	0.2	NaN	NaN	1.0	6.0	79.0	NaN	NaN	NaN	NaN	NaN	28
4	2014-06-01 01:00:00	NaN	NaN	NaN	NaN	1.0	6.0	75.0	NaN	NaN	4.0	NaN	NaN	28
...	...	...	...	...	...	...	...	...	...	...	...	...	...	...
210019	2014-09-01 00:00:00	NaN	0.5	NaN	NaN	20.0	84.0	29.0	NaN	NaN	NaN	NaN	NaN	28
210020	2014-09-01 00:00:00	NaN	0.3	NaN	NaN	1.0	22.0	NaN	15.0	NaN	6.0	NaN	NaN	28
210021	2014-09-01 00:00:00	NaN	NaN	NaN	NaN	1.0	13.0	70.0	NaN	NaN	NaN	NaN	NaN	28
210022	2014-09-01 00:00:00	NaN	NaN	NaN	NaN	3.0	38.0	42.0	NaN	NaN	NaN	NaN	NaN	28
210023	2014-09-01 00:00:00	NaN	NaN	NaN	NaN	1.0	26.0	65.0	11.0	NaN	NaN	NaN	NaN	28

210024 rows × 14 columns



```
In [3]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 210024 entries, 0 to 210023
Data columns (total 14 columns):
#   Column      Non-Null Count  Dtype
---  -
0   date        210024 non-null  object
1   BEN         46703 non-null   float64
2   CO          87023 non-null   float64
3   EBE         46722 non-null   float64
4   NMHC        25021 non-null   float64
5   NO          209154 non-null   float64
6   NO_2        209154 non-null   float64
7   O_3         121681 non-null   float64
8   PM10        104311 non-null   float64
9   PM25        51954 non-null   float64
10  SO_2        87141 non-null   float64
11  TCH         25021 non-null   float64
12  TOL         46570 non-null   float64
13  station     210024 non-null   int64
dtypes: float64(12), int64(1), object(1)
memory usage: 22.4+ MB
```

```
In [4]: df=df.dropna()
df
```

Out[4]:

	date	BEN	CO	EBE	NMHC	NO	NO_2	O_3	PM10	PM25	SO_2	TCH	TOL	s
<b>1</b>	2014-06-01 01:00:00	0.2	0.2	0.1	0.11	3.0	17.0	68.0	10.0	5.0	5.0	1.36	1.3	280
<b>6</b>	2014-06-01 01:00:00	0.1	0.2	0.1	0.23	1.0	5.0	80.0	4.0	3.0	2.0	1.21	0.1	280
<b>25</b>	2014-06-01 02:00:00	0.2	0.2	0.1	0.11	4.0	21.0	63.0	9.0	6.0	5.0	1.36	0.8	280
<b>30</b>	2014-06-01 02:00:00	0.2	0.2	0.1	0.23	1.0	4.0	88.0	7.0	5.0	2.0	1.21	0.1	280
<b>49</b>	2014-06-01 03:00:00	0.1	0.2	0.1	0.11	4.0	18.0	66.0	9.0	7.0	6.0	1.36	0.9	280
...	...	...	...	...	...	...	...	...	...	...	...	...	...	...
<b>209958</b>	2014-08-31 22:00:00	0.2	0.2	0.1	0.22	1.0	28.0	96.0	61.0	15.0	3.0	1.28	0.1	280
<b>209977</b>	2014-08-31 23:00:00	1.1	0.7	0.7	0.19	36.0	118.0	23.0	60.0	25.0	9.0	1.27	6.5	280
<b>209982</b>	2014-08-31 23:00:00	0.2	0.2	0.1	0.21	1.0	17.0	90.0	28.0	14.0	3.0	1.27	0.2	280
<b>210001</b>	2014-09-01 00:00:00	0.6	0.4	0.4	0.12	6.0	63.0	41.0	26.0	15.0	8.0	1.19	4.1	280
<b>210006</b>	2014-09-01 00:00:00	0.2	0.2	0.1	0.23	1.0	30.0	69.0	18.0	13.0	3.0	1.30	0.1	280

13946 rows × 14 columns



In [5]: `df.isnull().sum()`

```
Out[5]: date      0
      BEN      0
      CO      0
      EBE      0
      NMHC     0
      NO      0
      NO_2     0
      O_3      0
      PM10     0
      PM25     0
      SO_2     0
      TCH      0
      TOL      0
      station  0
      dtype: int64
```

In [6]: `df.describe()`

```
Out[6]:
```

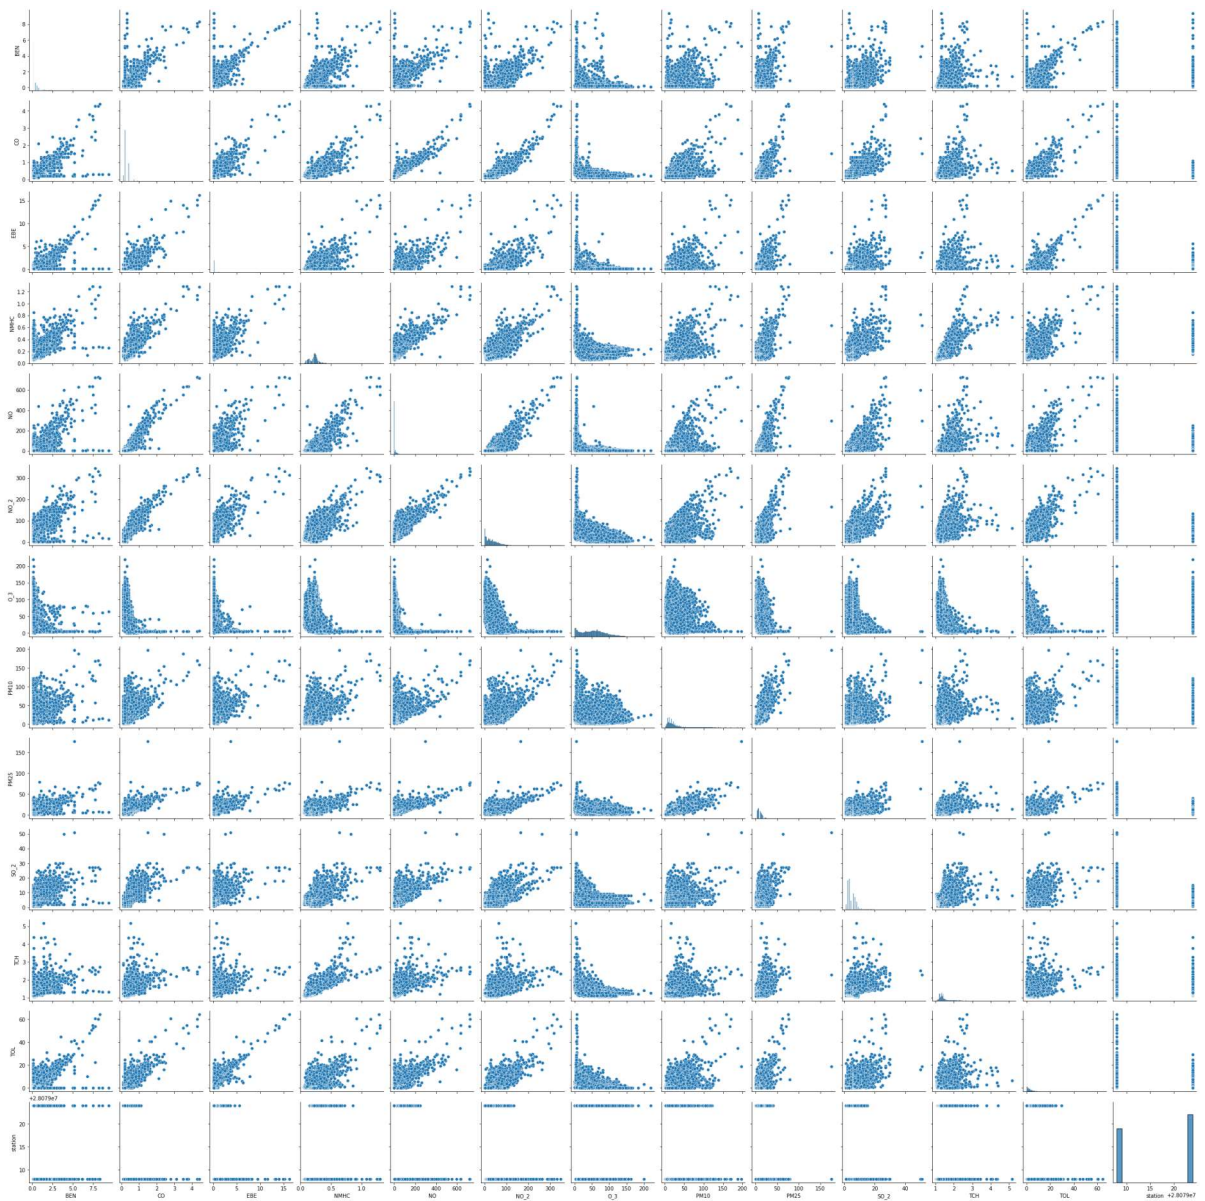
	BEN	CO	EBE	NMHC	NO	NO_2	
<b>count</b>	13946.000000	13946.000000	13946.000000	13946.000000	13946.000000	13946.000000	139
<b>mean</b>	0.375921	0.314793	0.306016	0.222302	17.589129	34.240929	
<b>std</b>	0.555093	0.207375	0.635475	0.082403	39.432216	30.654229	
<b>min</b>	0.100000	0.100000	0.100000	0.060000	1.000000	1.000000	
<b>25%</b>	0.100000	0.200000	0.100000	0.160000	1.000000	10.000000	
<b>50%</b>	0.200000	0.300000	0.100000	0.230000	4.000000	27.000000	
<b>75%</b>	0.400000	0.400000	0.300000	0.260000	18.000000	51.000000	
<b>max</b>	9.400000	4.400000	16.200001	1.290000	725.000000	346.000000	2

In [7]: `df.columns`

```
Out[7]: Index(['date', 'BEN', 'CO', 'EBE', 'NMHC', 'NO', 'NO_2', 'O_3', 'PM10', 'PM25',
              'SO_2', 'TCH', 'TOL', 'station'],
              dtype='object')
```

```
In [8]: sns.pairplot(df)
```

```
Out[8]: <seaborn.axisgrid.PairGrid at 0x288c9a79d60>
```




```
In [9]: sns.distplot(df['PM10'])
```

C:\ProgramData\Anaconda3\lib\site-packages\seaborn\distributions.py:2557: FutureWarning: `distplot` is a deprecated function and will be removed in a future version. Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).


```
warnings.warn(msg, FutureWarning)
```

```
Out[9]: <AxesSubplot:xlabel='PM10', ylabel='Density'>
```



```
In [10]: sns.heatmap(df.corr())
```

```
Out[10]: <AxesSubplot:>
```



```
In [11]: df.loc[df['TCH']<2,'TCH']=0
df.loc[df['TCH']>2,'TCH']=1
df['TCH']=df['TCH'].astype(int)
df
```

C:\ProgramData\Anaconda3\lib\site-packages\pandas\core\indexing.py:1720: SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame.

Try using .loc[row\_indexer,col\_indexer] = value instead

See the caveats in the documentation: [https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy) ([https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy))

```
self._setitem_single_column(loc, value, pi)
```

C:\ProgramData\Anaconda3\lib\site-packages\pandas\core\indexing.py:1720: SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame.

Try using .loc[row\_indexer,col\_indexer] = value instead

See the caveats in the documentation: [https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy) ([https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy))

```
self._setitem_single_column(loc, value, pi)
```

<ipython-input-11-e3d36a273982>:3: SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame.

Try using .loc[row\_indexer,col\_indexer] = value instead

See the caveats in the documentation: [https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy) ([https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy))

```
df['TCH']=df['TCH'].astype(int)
```



Out[11]:

	date	BEN	CO	EBE	NMHC	NO	NO_2	O_3	PM10	PM25	SO_2	TCH	TOL	station
1	2014-06-01 01:00:00	0.2	0.2	0.1	0.11	3.0	17.0	68.0	10.0	5.0	5.0	0	1.3	280
6	2014-06-01 01:00:00	0.1	0.2	0.1	0.23	1.0	5.0	80.0	4.0	3.0	2.0	0	0.1	280
25	2014-06-01 02:00:00	0.2	0.2	0.1	0.11	4.0	21.0	63.0	9.0	6.0	5.0	0	0.8	280
30	2014-06-01 02:00:00	0.2	0.2	0.1	0.23	1.0	4.0	88.0	7.0	5.0	2.0	0	0.1	280
49	2014-06-01 03:00:00	0.1	0.2	0.1	0.11	4.0	18.0	66.0	9.0	7.0	6.0	0	0.9	280
...	...	...	...	...	...	...	...	...	...	...	...	...	...	...
209958	2014-08-31 22:00:00	0.2	0.2	0.1	0.22	1.0	28.0	96.0	61.0	15.0	3.0	0	0.1	280
209977	2014-08-31 23:00:00	1.1	0.7	0.7	0.19	36.0	118.0	23.0	60.0	25.0	9.0	0	6.5	280
209982	2014-08-31 23:00:00	0.2	0.2	0.1	0.21	1.0	17.0	90.0	28.0	14.0	3.0	0	0.2	280
210001	2014-09-01 00:00:00	0.6	0.4	0.4	0.12	6.0	63.0	41.0	26.0	15.0	8.0	0	4.1	280
210006	2014-09-01 00:00:00	0.2	0.2	0.1	0.23	1.0	30.0	69.0	18.0	13.0	3.0	0	0.1	280

13946 rows × 14 columns



## LogisticRegression

```
In [12]: x=df[['BEN', 'CO', 'EBE', 'NMHC', 'NO', 'NO_2', 'O_3', 'PM10', 'PM25',
               'SO_2', 'TOL', 'station']]
y=df['TCH']
x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.3)
lgr=LogisticRegression()
lgr.fit(x_train,y_train)
```

Out[12]: LogisticRegression()

```
In [13]: lgr.predict(x_test)
```

```
Out[13]: array([0, 0, 0, ..., 0, 0, 0])
```

```
In [14]: lgr.score(x_test,y_test)
```

```
Out[14]: 0.9861376673040153
```

```
In [15]: fs=StandardScaler().fit_transform(x)
logr=LogisticRegression()
logr.fit(fs,y)
```

C:\ProgramData\Anaconda3\lib\site-packages\sklearn\linear\_model\\_logistic.py:  
763: ConvergenceWarning: lbfgs failed to converge (status=1):  
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max\_iter) or scale the data as shown in:  
<https://scikit-learn.org/stable/modules/preprocessing.html> (<https://scikit-learn.org/stable/modules/preprocessing.html>)  
Please also refer to the documentation for alternative solver options:  
[https://scikit-learn.org/stable/modules/linear\\_model.html#logistic-regression](https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression) ([https://scikit-learn.org/stable/modules/linear\\_model.html#logistic-regression](https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression))  
n\_iter\_i = \_check\_optimize\_result(

```
Out[15]: LogisticRegression()
```

```
In [16]: o=[[1,2,3,4,5,6,7,8,9,10,11,12]]
prediction=logr.predict(o)
print(prediction)
```

```
[0]
```

```
In [17]: logr.classes_
```

```
Out[17]: array([0, 1, 2])
```

```
In [18]: logr.predict_proba(o)[0][0]
```

```
Out[18]: 0.99999999999997202
```

```
In [19]: logr.predict_proba(o)[0][1]
```

```
Out[19]: 7.704648678920763e-15
```

## LinearRegression

```
In [20]: lr=LinearRegression()
lr.fit(x_train,y_train)
```

```
Out[20]: LinearRegression()
```

```
In [21]: print(lr.intercept_)
```

223968.26402332023


```
In [22]: coeff = pd.DataFrame(lr.coef_,x.columns,columns=['Co-efficient'])  
coeff
```

Out[22]:

	Co-efficient
<b>BEN</b>	-0.015361
<b>CO</b>	-0.207433
<b>EBE</b>	0.021501
<b>NMHC</b>	0.892985
<b>NO</b>	0.001524
<b>NO_2</b>	-0.000524
<b>O_3</b>	0.000280
<b>PM10</b>	-0.000286
<b>PM25</b>	0.002470
<b>SO_2</b>	-0.004989
<b>TOL</b>	-0.004160
<b>station</b>	-0.007976

```
In [23]: prediction=lr.predict(x_test)  
plt.scatter(y_test,prediction)
```

Out[23]: <matplotlib.collections.PathCollection at 0x288d959f340>



```
In [24]: print(lr.score(x_test,y_test))
```

0.322971710303267

## Ridge,Lasso

```
In [25]: rr=Ridge(alpha=10)
         rr.fit(x_train,y_train)
```

```
Out[25]: Ridge(alpha=10)
```

```
In [26]: rr.score(x_test,y_test)
```

```
Out[26]: 0.3053442466079368
```

```
In [27]: la=Lasso(alpha=10)
         la.fit(x_train,y_train)
```

```
Out[27]: Lasso(alpha=10)
```

```
In [28]: la.score(x_test,y_test)
```

```
Out[28]: -1.3761787409638515e-05
```

## ElasticNet

```
In [29]: en=ElasticNet()
         en.fit(x_train,y_train)
```

```
Out[29]: ElasticNet()
```

```
In [30]: print(en.coef_)
```

```
[0.         0.         0.         0.         0.00111633 0.
 0.         0.         0.         0.         0.         0.         ]
```

```
In [31]: print(en.intercept_)
```

```
-0.005319777935996429
```

```
In [32]: print(en.predict(x_train))
```

```
[-0.00308712  0.00026186  0.0270537  ...  0.01254145  0.00584349
 -0.00420345]
```

```
In [33]: print(en.score(x_train,y_train))
```

```
0.20531499203311165
```

```
In [34]: print("Mean Absolytre Error:",metrics.mean_absolute_error(y_test,prediction))
```

```
Mean Absolytre Error: 0.03859831113532434
```

```
In [35]: print("Mean Square Error:", metrics.mean_squared_error(y_test, prediction))
```

Mean Square Error: 0.009736006781566388

```
In [36]: print("Root Mean Square Error:", np.sqrt(metrics.mean_absolute_error(y_test, prediction)))
```

Root Mean Square Error: 0.19646452894943742

## RandomForest

```
In [37]: rfc=RandomForestClassifier()  
rfc.fit(x_train,y_train)
```

Out[37]: RandomForestClassifier()

```
In [38]: parameters={'max_depth':[1,2,3,4,5],  
                    'min_samples_leaf':[5,10,15,20,25],  
                    'n_estimators':[10,20,30,40,50]}
```

```
In [39]: grid_search=GridSearchCV(estimator=rfc,param_grid=parameters,cv=2,scoring="accuracy")  
grid_search.fit(x_train,y_train)
```

Out[39]: GridSearchCV(cv=2, estimator=RandomForestClassifier(),  
param\_grid={'max\_depth': [1, 2, 3, 4, 5],  
            'min\_samples\_leaf': [5, 10, 15, 20, 25],  
            'n\_estimators': [10, 20, 30, 40, 50]},  
scoring='accuracy')

```
In [40]: grid_search.best_score_
```

Out[40]: 0.9905757017004713

```
In [41]: rfc_best=grid_search.best_estimator_
```

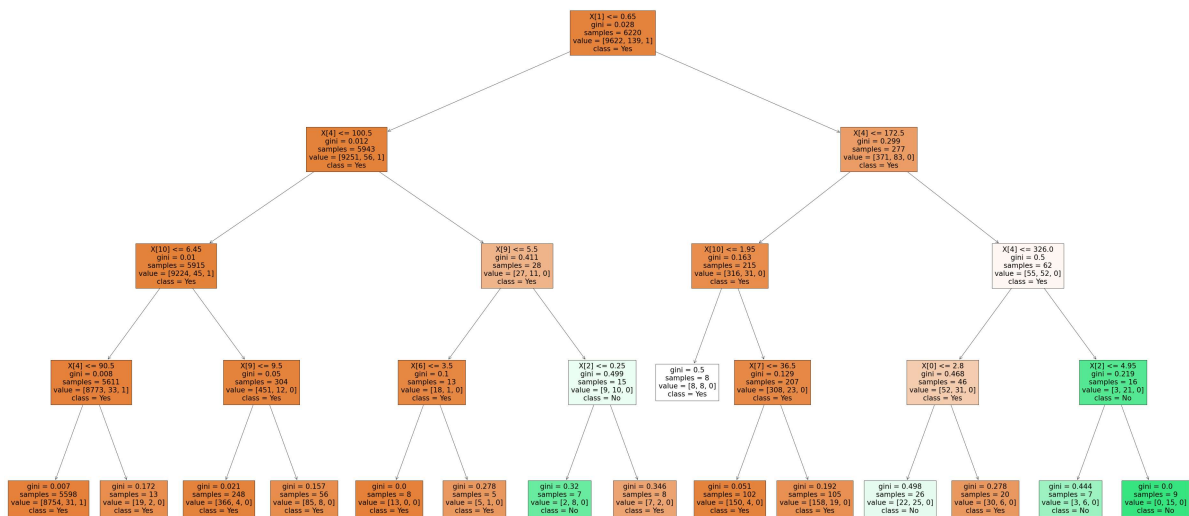
```
In [42]: plt.figure(figsize=(80,40))  
plot_tree(rfc_best.estimators_[5],class_names=['Yes','No','Yes','No'],filled=True)
```

```

Out[42]: [Text(2271.8571428571427, 1956.96, 'X[1] <= 0.65\ngini = 0.028\nsamples = 622\nvalue = [9622, 139, 1]\nclass = Yes'),
Text(1275.4285714285713, 1522.0800000000002, 'X[4] <= 100.5\ngini = 0.012\nsamples = 5943\nvalue = [9251, 56, 1]\nclass = Yes'),
Text(637.7142857142857, 1087.2, 'X[10] <= 6.45\ngini = 0.01\nsamples = 5915\nvalue = [9224, 45, 1]\nclass = Yes'),
Text(318.85714285714283, 652.3200000000002, 'X[4] <= 90.5\ngini = 0.008\nsamples = 5611\nvalue = [8773, 33, 1]\nclass = Yes'),
Text(159.42857142857142, 217.44000000000005, 'gini = 0.007\nsamples = 5598\nvalue = [8754, 31, 1]\nclass = Yes'),
Text(478.2857142857142, 217.44000000000005, 'gini = 0.172\nsamples = 13\nvalue = [19, 2, 0]\nclass = Yes'),
Text(956.5714285714284, 652.3200000000002, 'X[9] <= 9.5\ngini = 0.05\nsamples = 304\nvalue = [451, 12, 0]\nclass = Yes'),
Text(797.1428571428571, 217.44000000000005, 'gini = 0.021\nsamples = 248\nvalue = [366, 4, 0]\nclass = Yes'),
Text(1116.0, 217.44000000000005, 'gini = 0.157\nsamples = 56\nvalue = [85, 8, 0]\nclass = Yes'),
Text(1913.1428571428569, 1087.2, 'X[9] <= 5.5\ngini = 0.411\nsamples = 28\nvalue = [27, 11, 0]\nclass = Yes'),
Text(1594.2857142857142, 652.3200000000002, 'X[6] <= 3.5\ngini = 0.1\nsamples = 13\nvalue = [18, 1, 0]\nclass = Yes'),
Text(1434.8571428571427, 217.44000000000005, 'gini = 0.0\nsamples = 8\nvalue = [13, 0, 0]\nclass = Yes'),
Text(1753.7142857142856, 217.44000000000005, 'gini = 0.278\nsamples = 5\nvalue = [5, 1, 0]\nclass = Yes'),
Text(2232.0, 652.3200000000002, 'X[2] <= 0.25\ngini = 0.499\nsamples = 15\nvalue = [9, 10, 0]\nclass = No'),
Text(2072.5714285714284, 217.44000000000005, 'gini = 0.32\nsamples = 7\nvalue = [2, 8, 0]\nclass = No'),
Text(2391.428571428571, 217.44000000000005, 'gini = 0.346\nsamples = 8\nvalue = [7, 2, 0]\nclass = Yes'),
Text(3268.285714285714, 1522.0800000000002, 'X[4] <= 172.5\ngini = 0.299\nsamples = 277\nvalue = [371, 83, 0]\nclass = Yes'),
Text(2710.285714285714, 1087.2, 'X[10] <= 1.95\ngini = 0.163\nsamples = 215\nvalue = [316, 31, 0]\nclass = Yes'),
Text(2550.8571428571427, 652.3200000000002, 'gini = 0.5\nsamples = 8\nvalue = [8, 8, 0]\nclass = Yes'),
Text(2869.7142857142853, 652.3200000000002, 'X[7] <= 36.5\ngini = 0.129\nsamples = 207\nvalue = [308, 23, 0]\nclass = Yes'),
Text(2710.285714285714, 217.44000000000005, 'gini = 0.051\nsamples = 102\nvalue = [150, 4, 0]\nclass = Yes'),
Text(3029.142857142857, 217.44000000000005, 'gini = 0.192\nsamples = 105\nvalue = [158, 19, 0]\nclass = Yes'),
Text(3826.2857142857138, 1087.2, 'X[4] <= 326.0\ngini = 0.5\nsamples = 62\nvalue = [55, 52, 0]\nclass = Yes'),
Text(3507.428571428571, 652.3200000000002, 'X[0] <= 2.8\ngini = 0.468\nsamples = 46\nvalue = [52, 31, 0]\nclass = Yes'),
Text(3347.9999999999995, 217.44000000000005, 'gini = 0.498\nsamples = 26\nvalue = [22, 25, 0]\nclass = No'),
Text(3666.8571428571427, 217.44000000000005, 'gini = 0.278\nsamples = 20\nvalue = [30, 6, 0]\nclass = Yes'),
Text(4145.142857142857, 652.3200000000002, 'X[2] <= 4.95\ngini = 0.219\nsamples = 16\nvalue = [3, 21, 0]\nclass = No'),
Text(3985.7142857142853, 217.44000000000005, 'gini = 0.444\nsamples = 7\nvalue = [3, 6, 0]\nclass = No'),

```

```
Text(4304.571428571428, 217.44000000000005, 'gini = 0.0\nsamples = 9\nvalue  
= [0, 15, 0]\nnclass = No')]
```



Best model:LogisticRegression

In [ ]: