

```
In [1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.linear_model import Ridge,Lasso
from sklearn.linear_model import ElasticNet
from sklearn import metrics
from sklearn.linear_model import LogisticRegression
from sklearn.preprocessing import StandardScaler
from sklearn.ensemble import RandomForestClassifier
from sklearn.model_selection import GridSearchCV
from sklearn.tree import plot_tree
```

```
In [2]: df=pd.read_csv(r"C:\Users\user\Downloads\csvs_per_year\csvs_per_year\madrid_2015.csv")
df
```

Out[2]:

	date	BEN	CO	EBE	NMHC	NO	NO_2	O_3	PM10	PM25	SO_2	TCH	TOL	
0	2015-10-01 01:00:00	NaN	0.8	NaN	NaN	90.0	82.0	NaN	NaN	NaN	10.0	NaN	NaN	28
1	2015-10-01 01:00:00	2.0	0.8	1.6	0.33	40.0	95.0	4.0	37.0	24.0	12.0	1.83	8.3	28
2	2015-10-01 01:00:00	3.1	NaN	1.8	NaN	29.0	97.0	NaN	NaN	NaN	NaN	NaN	7.1	28
3	2015-10-01 01:00:00	NaN	0.6	NaN	NaN	30.0	103.0	2.0	NaN	NaN	NaN	NaN	NaN	28
4	2015-10-01 01:00:00	NaN	NaN	NaN	NaN	95.0	96.0	2.0	NaN	NaN	9.0	NaN	NaN	28
...	...	...	...	...	...	...	...	...	...	...	...	...	...	...
210091	2015-08-01 00:00:00	NaN	0.2	NaN	NaN	11.0	33.0	53.0	NaN	NaN	NaN	NaN	NaN	28
210092	2015-08-01 00:00:00	NaN	0.2	NaN	NaN	1.0	5.0	NaN	26.0	NaN	10.0	NaN	NaN	28
210093	2015-08-01 00:00:00	NaN	NaN	NaN	NaN	1.0	7.0	74.0	NaN	NaN	NaN	NaN	NaN	28
210094	2015-08-01 00:00:00	NaN	NaN	NaN	NaN	3.0	7.0	65.0	NaN	NaN	NaN	NaN	NaN	28
210095	2015-08-01 00:00:00	NaN	NaN	NaN	NaN	1.0	9.0	54.0	29.0	NaN	NaN	NaN	NaN	28

210096 rows × 14 columns



```
In [3]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 210096 entries, 0 to 210095
Data columns (total 14 columns):
 #   Column      Non-Null Count  Dtype
---  -
 0   date        210096 non-null object
 1   BEN         51039 non-null float64
 2   CO          86827 non-null float64
 3   EBE         50962 non-null float64
 4   NMHC        25756 non-null float64
 5   NO          208805 non-null float64
 6   NO_2        208805 non-null float64
 7   O_3         121574 non-null float64
 8   PM10        102745 non-null float64
 9   PM25        48798 non-null float64
10   SO_2        86898 non-null float64
11   TCH         25756 non-null float64
12   TOL         50626 non-null float64
13   station     210096 non-null int64
dtypes: float64(12), int64(1), object(1)
memory usage: 22.4+ MB
```

```
In [4]: df=df.dropna()
df
```

Out[4]:

	date	BEN	CO	EBE	NMHC	NO	NO_2	O_3	PM10	PM25	SO_2	TCH	TOL	
<b>1</b>	2015-10-01 01:00:00	2.0	0.8	1.6	0.33	40.0	95.0	4.0	37.0	24.0	12.0	1.83	8.3	28
<b>6</b>	2015-10-01 01:00:00	0.5	0.3	0.3	0.12	6.0	83.0	1.0	19.0	12.0	3.0	1.29	4.8	28
<b>25</b>	2015-10-01 02:00:00	1.6	0.7	1.3	0.38	81.0	105.0	4.0	36.0	19.0	13.0	1.93	6.9	28
<b>30</b>	2015-10-01 02:00:00	0.4	0.3	0.3	0.11	5.0	72.0	2.0	16.0	10.0	2.0	1.27	7.8	28
<b>49</b>	2015-10-01 03:00:00	2.2	0.8	1.8	0.41	111.0	104.0	4.0	35.0	20.0	14.0	2.05	13.9	28
...	...	...	...	...	...	...	...	...	...	...	...	...	...	...
<b>210030</b>	2015-07-31 22:00:00	0.1	0.1	0.1	0.06	1.0	10.0	69.0	10.0	3.0	2.0	1.18	0.2	28
<b>210049</b>	2015-07-31 23:00:00	0.4	0.3	0.1	0.12	3.0	28.0	56.0	15.0	7.0	12.0	1.45	1.2	28
<b>210054</b>	2015-07-31 23:00:00	0.1	0.1	0.1	0.06	1.0	10.0	63.0	5.0	1.0	2.0	1.18	0.2	28
<b>210073</b>	2015-08-01 00:00:00	0.1	0.3	0.1	0.11	2.0	23.0	59.0	5.0	2.0	11.0	1.44	0.6	28
<b>210078</b>	2015-08-01 00:00:00	0.1	0.1	0.1	0.06	1.0	8.0	65.0	7.0	1.0	2.0	1.18	0.4	28

16026 rows × 14 columns



In [5]: `df.isnull().sum()`

```
Out[5]: date      0
      BEN      0
      CO      0
      EBE      0
      NMHC     0
      NO      0
      NO_2     0
      O_3      0
      PM10     0
      PM25     0
      SO_2     0
      TCH      0
      TOL      0
      station  0
      dtype: int64
```

In [6]: `df.describe()`

```
Out[6]:
```

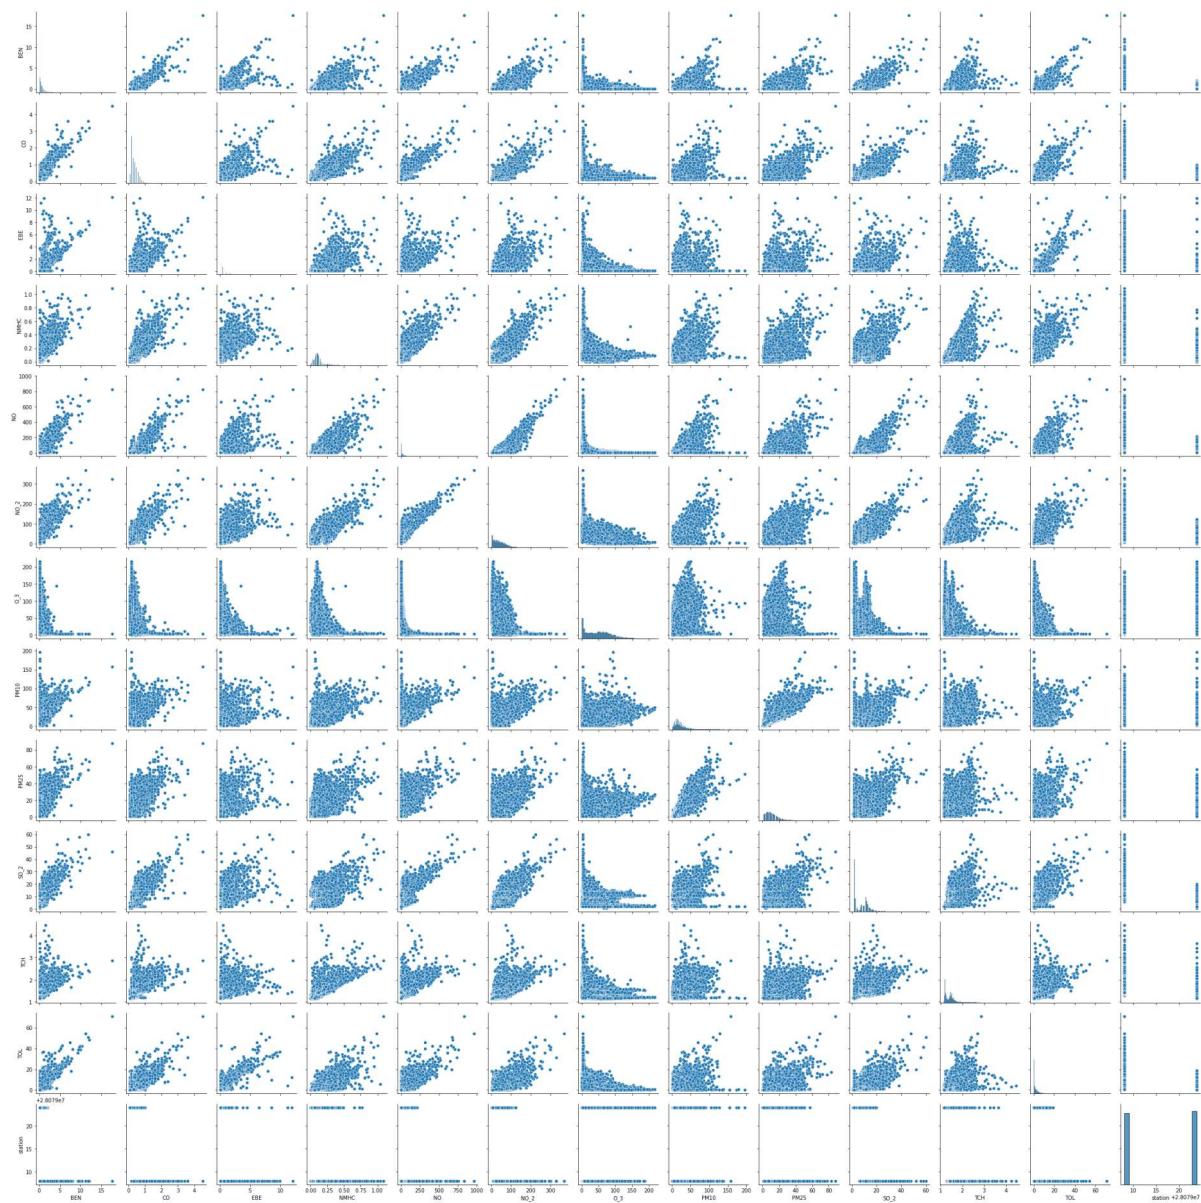
	BEN	CO	EBE	NMHC	NO	NO_2	
<b>count</b>	16026.000000	16026.000000	16026.000000	16026.000000	16026.000000	16026.000000	160
<b>mean</b>	0.504823	0.380594	0.394247	0.123099	23.842256	40.948771	
<b>std</b>	0.716896	0.260805	0.678592	0.092368	51.255660	33.236098	
<b>min</b>	0.100000	0.100000	0.100000	0.000000	1.000000	1.000000	
<b>25%</b>	0.100000	0.200000	0.100000	0.070000	1.000000	14.000000	
<b>50%</b>	0.200000	0.300000	0.100000	0.100000	6.000000	35.000000	
<b>75%</b>	0.700000	0.500000	0.400000	0.140000	24.000000	60.000000	
<b>max</b>	17.700001	4.500000	12.100000	1.090000	960.000000	369.000000	2

In [7]: `df.columns`

```
Out[7]: Index(['date', 'BEN', 'CO', 'EBE', 'NMHC', 'NO', 'NO_2', 'O_3', 'PM10', 'PM25',
              'SO_2', 'TCH', 'TOL', 'station'],
              dtype='object')
```

```
In [8]: sns.pairplot(df)
```

```
Out[8]: <seaborn.axisgrid.PairGrid at 0x1a5138711c0>
```




```
In [9]: sns.distplot(df['PM10'])
```

C:\ProgramData\Anaconda3\lib\site-packages\seaborn\distributions.py:2557: FutureWarning: `distplot` is a deprecated function and will be removed in a future version. Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).


```
warnings.warn(msg, FutureWarning)
```

```
Out[9]: <AxesSubplot:xlabel='PM10', ylabel='Density'>
```



```
In [10]: sns.heatmap(df.corr())
```

```
Out[10]: <AxesSubplot:>
```



```
In [11]: df.loc[df['TCH']<2,'TCH']=0
df.loc[df['TCH']>2,'TCH']=1
df['TCH']=df['TCH'].astype(int)
df
```

C:\ProgramData\Anaconda3\lib\site-packages\pandas\core\indexing.py:1720: SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame.

Try using .loc[row\_indexer,col\_indexer] = value instead

See the caveats in the documentation: [https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy) ([https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy))

```
self._setitem_single_column(loc, value, pi)
```

C:\ProgramData\Anaconda3\lib\site-packages\pandas\core\indexing.py:1720: SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame.

Try using .loc[row\_indexer,col\_indexer] = value instead

See the caveats in the documentation: [https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy) ([https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy))

```
self._setitem_single_column(loc, value, pi)
```

<ipython-input-11-e3d36a273982>:3: SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame.

Try using .loc[row\_indexer,col\_indexer] = value instead

See the caveats in the documentation: [https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy) ([https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy))

```
df['TCH']=df['TCH'].astype(int)
```



Out[11]:

	date	BEN	CO	EBE	NMHC	NO	NO_2	O_3	PM10	PM25	SO_2	TCH	TOL
1	2015-10-01 01:00:00	2.0	0.8	1.6	0.33	40.0	95.0	4.0	37.0	24.0	12.0	0	8.3 28
6	2015-10-01 01:00:00	0.5	0.3	0.3	0.12	6.0	83.0	1.0	19.0	12.0	3.0	0	4.8 28
25	2015-10-01 02:00:00	1.6	0.7	1.3	0.38	81.0	105.0	4.0	36.0	19.0	13.0	0	6.9 28
30	2015-10-01 02:00:00	0.4	0.3	0.3	0.11	5.0	72.0	2.0	16.0	10.0	2.0	0	7.8 28
49	2015-10-01 03:00:00	2.2	0.8	1.8	0.41	111.0	104.0	4.0	35.0	20.0	14.0	1	13.9 28
...	...	...	...	...	...	...	...	...	...	...	...	...	...
210030	2015-07-31 22:00:00	0.1	0.1	0.1	0.06	1.0	10.0	69.0	10.0	3.0	2.0	0	0.2 28
210049	2015-07-31 23:00:00	0.4	0.3	0.1	0.12	3.0	28.0	56.0	15.0	7.0	12.0	0	1.2 28
210054	2015-07-31 23:00:00	0.1	0.1	0.1	0.06	1.0	10.0	63.0	5.0	1.0	2.0	0	0.2 28
210073	2015-08-01 00:00:00	0.1	0.3	0.1	0.11	2.0	23.0	59.0	5.0	2.0	11.0	0	0.6 28
210078	2015-08-01 00:00:00	0.1	0.1	0.1	0.06	1.0	8.0	65.0	7.0	1.0	2.0	0	0.4 28

16026 rows × 14 columns



## LogisticRegression

```
In [12]: x=df[['BEN', 'CO', 'EBE', 'NMHC', 'NO', 'NO_2', 'O_3', 'PM10', 'PM25',
               'SO_2', 'TOL', 'station']]
y=df['TCH']
x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.3)
lgr=LogisticRegression()
lgr.fit(x_train,y_train)
```

Out[12]: LogisticRegression()

```
In [13]: lgr.predict(x_test)
```

```
Out[13]: array([0, 0, 0, ..., 0, 0, 0])
```

```
In [14]: lgr.score(x_test,y_test)
```

```
Out[14]: 0.9762895174708819
```

```
In [15]: fs=StandardScaler().fit_transform(x)
logr=LogisticRegression()
logr.fit(fs,y)
```

C:\ProgramData\Anaconda3\lib\site-packages\sklearn\linear\_model\\_logistic.py:  
763: ConvergenceWarning: lbfgs failed to converge (status=1):  
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max\_iter) or scale the data as shown in:  
<https://scikit-learn.org/stable/modules/preprocessing.html> (<https://scikit-learn.org/stable/modules/preprocessing.html>)  
Please also refer to the documentation for alternative solver options:  
[https://scikit-learn.org/stable/modules/linear\\_model.html#logistic-regression](https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression) ([https://scikit-learn.org/stable/modules/linear\\_model.html#logistic-regression](https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression))  
n\_iter\_i = \_check\_optimize\_result(

```
Out[15]: LogisticRegression()
```

```
In [16]: o=[[1,2,3,4,5,6,7,8,9,10,11,12]]
prediction=logr.predict(o)
print(prediction)
```

```
[0]
```

```
In [17]: logr.classes_
```

```
Out[17]: array([0, 1, 2])
```

```
In [18]: logr.predict_proba(o)[0][0]
```

```
Out[18]: 0.9999196786297907
```

```
In [19]: logr.predict_proba(o)[0][1]
```

```
Out[19]: 2.053888893534883e-14
```

## LinearRegression

```
In [20]: lr=LinearRegression()
lr.fit(x_train,y_train)
```

```
Out[20]: LinearRegression()
```

```
In [21]: print(lr.intercept_)
```

52904.227763502706


```
In [22]: coeff = pd.DataFrame(lr.coef_,x.columns,columns=['Co-efficient'])  
coeff
```

Out[22]:

	Co-efficient
<b>BEN</b>	0.027972
<b>CO</b>	-0.074742
<b>EBE</b>	0.006364
<b>NMHC</b>	0.572787
<b>NO</b>	0.001284
<b>NO_2</b>	-0.001044
<b>O_3</b>	0.000014
<b>PM10</b>	0.000890
<b>PM25</b>	-0.000281
<b>SO_2</b>	-0.000795
<b>TOL</b>	-0.004117
<b>station</b>	-0.001884

```
In [23]: prediction=lr.predict(x_test)  
plt.scatter(y_test,prediction)
```

Out[23]: <matplotlib.collections.PathCollection at 0x1a5238feb50>



```
In [24]: print(lr.score(x_test,y_test))
```

0.2556798222158476

## Ridge,Lasso

```
In [25]: rr=Ridge(alpha=10)
         rr.fit(x_train,y_train)
```

```
Out[25]: Ridge(alpha=10)
```

```
In [26]: rr.score(x_test,y_test)
```

```
Out[26]: 0.2559853179539364
```

```
In [27]: la=Lasso(alpha=10)
         la.fit(x_train,y_train)
```

```
Out[27]: Lasso(alpha=10)
```

```
In [28]: la.score(x_test,y_test)
```

```
Out[28]: -3.4146404641610673e-06
```

## ElasticNet

```
In [29]: en=ElasticNet()
         en.fit(x_train,y_train)
```

```
Out[29]: ElasticNet()
```

```
In [30]: print(en.coef_)
```

```
[ 0.          0.          0.          0.          0.00127155  0.
  0.          0.          0.          0.          0.         -0.         ]
```

```
In [31]: print(en.intercept_)
```

```
-0.006090643490251853
```

```
In [32]: print(en.predict(x_train))
```

```
[-0.002276    0.06002987  0.02442652 ...  0.06257297  0.14268052
  0.00535329]
```

```
In [33]: print(en.score(x_train,y_train))
```

```
0.22078010833943218
```

```
In [34]: print("Mean Absolytre Error:",metrics.mean_absolute_error(y_test,prediction))
```

```
Mean Absolytre Error: 0.05221831972592315
```

```
In [35]: print("Mean Square Error:", metrics.mean_squared_error(y_test, prediction))
```

Mean Square Error: 0.019057567847766688

```
In [36]: print("Root Mean Square Error:", np.sqrt(metrics.mean_absolute_error(y_test, prediction)))
```

Root Mean Square Error: 0.22851328129000106

## RandomForest

```
In [37]: rfc=RandomForestClassifier()  
rfc.fit(x_train,y_train)
```

Out[37]: RandomForestClassifier()

```
In [38]: parameters={'max_depth':[1,2,3,4,5],  
                    'min_samples_leaf':[5,10,15,20,25],  
                    'n_estimators':[10,20,30,40,50]}
```

```
In [39]: grid_search=GridSearchCV(estimator=rfc,param_grid=parameters,cv=2,scoring="accuracy")  
grid_search.fit(x_train,y_train)
```

Out[39]: GridSearchCV(cv=2, estimator=RandomForestClassifier(),  
param\_grid={'max\_depth': [1, 2, 3, 4, 5],  
 'min\_samples\_leaf': [5, 10, 15, 20, 25],  
 'n\_estimators': [10, 20, 30, 40, 50]},  
scoring='accuracy')

```
In [40]: grid_search.best_score_
```

Out[40]: 0.9841326439650562

```
In [41]: rfc_best=grid_search.best_estimator_
```

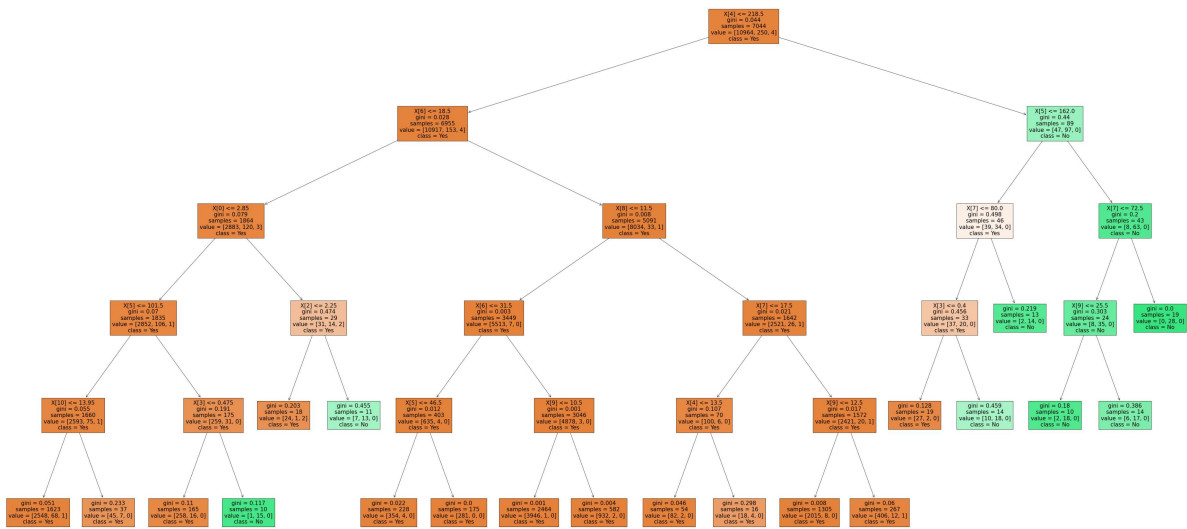
```
In [42]: plt.figure(figsize=(80,40))  
plot_tree(rfc_best.estimators_[5],class_names=['Yes','No','Yes','No'],filled=True)
```

```

Out[42]: [Text(2773.5882352941176, 1993.2, 'X[4] <= 218.5\ngini = 0.044\nsamples = 704
4\nvalue = [10964, 250, 4]\nclclass = Yes'),
Text(1608.3529411764705, 1630.8000000000002, 'X[6] <= 18.5\ngini = 0.028\nsa
mples = 6955\nvalue = [10917, 153, 4]\nclclass = Yes'),
Text(853.4117647058823, 1268.4, 'X[0] <= 2.85\ngini = 0.079\nsamples = 1864
\nvalue = [2883, 120, 3]\nclclass = Yes'),
Text(525.1764705882352, 906.0, 'X[5] <= 101.5\ngini = 0.07\nsamples = 1835\n
value = [2852, 106, 1]\nclclass = Yes'),
Text(262.5882352941176, 543.5999999999999, 'X[10] <= 13.95\ngini = 0.055\nsa
mples = 1660\nvalue = [2593, 75, 1]\nclclass = Yes'),
Text(131.2941176470588, 181.1999999999998, 'gini = 0.051\nsamples = 1623\nv
alue = [2548, 68, 1]\nclclass = Yes'),
Text(393.88235294117646, 181.1999999999998, 'gini = 0.233\nsamples = 37\nva
lue = [45, 7, 0]\nclclass = Yes'),
Text(787.7647058823529, 543.5999999999999, 'X[3] <= 0.475\ngini = 0.191\nsam
ples = 175\nvalue = [259, 31, 0]\nclclass = Yes'),
Text(656.470588235294, 181.1999999999998, 'gini = 0.11\nsamples = 165\nvalu
e = [258, 16, 0]\nclclass = Yes'),
Text(919.0588235294117, 181.1999999999998, 'gini = 0.117\nsamples = 10\nval
ue = [1, 15, 0]\nclclass = No'),
Text(1181.6470588235293, 906.0, 'X[2] <= 2.25\ngini = 0.474\nsamples = 29\nv
alue = [31, 14, 2]\nclclass = Yes'),
Text(1050.3529411764705, 543.5999999999999, 'gini = 0.203\nsamples = 18\nval
ue = [24, 1, 2]\nclclass = Yes'),
Text(1312.941176470588, 543.5999999999999, 'gini = 0.455\nsamples = 11\nvalu
e = [7, 13, 0]\nclclass = No'),
Text(2363.2941176470586, 1268.4, 'X[8] <= 11.5\ngini = 0.008\nsamples = 5091
\nvalue = [8034, 33, 1]\nclclass = Yes'),
Text(1838.1176470588234, 906.0, 'X[6] <= 31.5\ngini = 0.003\nsamples = 3449
\nvalue = [5513, 7, 0]\nclclass = Yes'),
Text(1575.5294117647059, 543.5999999999999, 'X[5] <= 46.5\ngini = 0.012\nsam
ples = 403\nvalue = [635, 4, 0]\nclclass = Yes'),
Text(1444.2352941176468, 181.1999999999998, 'gini = 0.022\nsamples = 228\nv
alue = [354, 4, 0]\nclclass = Yes'),
Text(1706.8235294117646, 181.1999999999998, 'gini = 0.0\nsamples = 175\nval
ue = [281, 0, 0]\nclclass = Yes'),
Text(2100.705882352941, 543.5999999999999, 'X[9] <= 10.5\ngini = 0.001\nsamp
les = 3046\nvalue = [4878, 3, 0]\nclclass = Yes'),
Text(1969.4117647058822, 181.1999999999998, 'gini = 0.001\nsamples = 2464\n
value = [3946, 1, 0]\nclclass = Yes'),
Text(2232.0, 181.1999999999998, 'gini = 0.004\nsamples = 582\nvalue = [932,
2, 0]\nclclass = Yes'),
Text(2888.4705882352937, 906.0, 'X[7] <= 17.5\ngini = 0.021\nsamples = 1642
\nvalue = [2521, 26, 1]\nclclass = Yes'),
Text(2625.882352941176, 543.5999999999999, 'X[4] <= 13.5\ngini = 0.107\nsamp
les = 70\nvalue = [100, 6, 0]\nclclass = Yes'),
Text(2494.5882352941176, 181.1999999999998, 'gini = 0.046\nsamples = 54\nva
lue = [82, 2, 0]\nclclass = Yes'),
Text(2757.176470588235, 181.1999999999998, 'gini = 0.298\nsamples = 16\nval
ue = [18, 4, 0]\nclclass = Yes'),
Text(3151.0588235294117, 543.5999999999999, 'X[9] <= 12.5\ngini = 0.017\nsam
ples = 1572\nvalue = [2421, 20, 1]\nclclass = Yes'),
Text(3019.7647058823527, 181.1999999999998, 'gini = 0.008\nsamples = 1305\n
value = [2015, 8, 0]\nclclass = Yes'),
Text(3282.3529411764703, 181.1999999999998, 'gini = 0.06\nsamples = 267\nva
lue = [406, 12, 1]\nclclass = Yes'),
Text(3938.8235294117644, 1630.8000000000002, 'X[5] <= 162.0\ngini = 0.44\nsa

```

```
mples = 89\nvalue = [47, 97, 0]\nnclass = No'),  
  Text(3676.235294117647, 1268.4, 'X[7] <= 80.0\ngini = 0.498\nsamples = 46\nvalue = [39, 34, 0]\nnclass = Yes'),  
  Text(3544.941176470588, 906.0, 'X[3] <= 0.4\ngini = 0.456\nsamples = 33\nvalue = [37, 20, 0]\nnclass = Yes'),  
  Text(3413.6470588235293, 543.5999999999999, 'gini = 0.128\nsamples = 19\nvalue = [27, 2, 0]\nnclass = Yes'),  
  Text(3676.235294117647, 543.5999999999999, 'gini = 0.459\nsamples = 14\nvalue = [10, 18, 0]\nnclass = No'),  
  Text(3807.5294117647054, 906.0, 'gini = 0.219\nsamples = 13\nvalue = [2, 14, 0]\nnclass = No'),  
  Text(4201.411764705882, 1268.4, 'X[7] <= 72.5\ngini = 0.2\nsamples = 43\nvalue = [8, 63, 0]\nnclass = No'),  
  Text(4070.117647058823, 906.0, 'X[9] <= 25.5\ngini = 0.303\nsamples = 24\nvalue = [8, 35, 0]\nnclass = No'),  
  Text(3938.8235294117644, 543.5999999999999, 'gini = 0.18\nsamples = 10\nvalue = [2, 18, 0]\nnclass = No'),  
  Text(4201.411764705882, 543.5999999999999, 'gini = 0.386\nsamples = 14\nvalue = [6, 17, 0]\nnclass = No'),  
  Text(4332.7058823529405, 906.0, 'gini = 0.0\nsamples = 19\nvalue = [0, 28, 0]\nnclass = No')]
```



Best model:LogisticRegression

```
In [ ]:
```