

```
In [1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.linear_model import Ridge,Lasso
from sklearn.linear_model import ElasticNet
from sklearn import metrics
from sklearn.linear_model import LogisticRegression
from sklearn.preprocessing import StandardScaler
from sklearn.ensemble import RandomForestClassifier
from sklearn.model_selection import GridSearchCV
from sklearn.tree import plot_tree
```

```
In [2]: df=pd.read_csv(r"C:\Users\user\Downloads\csvs_per_year\csvs_per_year\madrid_2016\
df
```

Out[2]:

	date	BEN	CO	EBE	NMHC	NO	NO_2	O_3	PM10	PM25	SO_2	TCH	TOL	
0	2016-11-01 01:00:00	NaN	0.7	NaN	NaN	153.0	77.0	NaN	NaN	NaN	7.0	NaN	NaN	2
1	2016-11-01 01:00:00	3.1	1.1	2.0	0.53	260.0	144.0	4.0	46.0	24.0	18.0	2.44	14.4	2
2	2016-11-01 01:00:00	5.9	NaN	7.5	NaN	297.0	139.0	NaN	NaN	NaN	NaN	NaN	26.0	2
3	2016-11-01 01:00:00	NaN	1.0	NaN	NaN	154.0	113.0	2.0	NaN	NaN	NaN	NaN	NaN	2
4	2016-11-01 01:00:00	NaN	NaN	NaN	NaN	275.0	127.0	2.0	NaN	NaN	18.0	NaN	NaN	2
...	
209491	2016-07-01 00:00:00	NaN	0.2	NaN	NaN	2.0	29.0	73.0	NaN	NaN	NaN	NaN	NaN	2
209492	2016-07-01 00:00:00	NaN	0.3	NaN	NaN	1.0	29.0	NaN	36.0	NaN	5.0	NaN	NaN	2
209493	2016-07-01 00:00:00	NaN	NaN	NaN	NaN	1.0	19.0	71.0	NaN	NaN	NaN	NaN	NaN	2
209494	2016-07-01 00:00:00	NaN	NaN	NaN	NaN	6.0	17.0	85.0	NaN	NaN	NaN	NaN	NaN	2
209495	2016-07-01 00:00:00	NaN	NaN	NaN	NaN	2.0	46.0	61.0	34.0	NaN	NaN	NaN	NaN	2

209496 rows × 14 columns



```
In [3]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 209496 entries, 0 to 209495
Data columns (total 14 columns):
#   Column      Non-Null Count  Dtype
---  -
0   date        209496 non-null object
1   BEN         50755 non-null  float64
2   CO          85999 non-null  float64
3   EBE         50335 non-null  float64
4   NMHC        25970 non-null  float64
5   NO          208614 non-null float64
6   NO_2        208614 non-null float64
7   O_3         121197 non-null float64
8   PM10        102892 non-null float64
9   PM25        52165 non-null  float64
10  SO_2        86023 non-null  float64
11  TCH         25970 non-null  float64
12  TOL         50662 non-null  float64
13  station     209496 non-null int64
dtypes: float64(12), int64(1), object(1)
memory usage: 22.4+ MB
```

```
In [4]: df=df.dropna()
df
```

Out[4]:

	date	BEN	CO	EBE	NMHC	NO	NO_2	O_3	PM10	PM25	SO_2	TCH	TOL	
1	2016-11-01 01:00:00	3.1	1.1	2.0	0.53	260.0	144.0	4.0	46.0	24.0	18.0	2.44	14.4	28
6	2016-11-01 01:00:00	0.7	0.8	0.4	0.13	57.0	66.0	3.0	23.0	15.0	4.0	1.35	5.0	28
25	2016-11-01 02:00:00	2.7	1.0	2.1	0.40	139.0	114.0	4.0	37.0	21.0	14.0	2.30	15.0	28
30	2016-11-01 02:00:00	0.7	0.7	0.4	0.13	48.0	59.0	3.0	23.0	15.0	3.0	1.35	5.0	28
49	2016-11-01 03:00:00	1.7	0.8	1.4	0.25	53.0	90.0	4.0	31.0	19.0	10.0	1.95	10.7	28
...
209430	2016-06-30 22:00:00	0.1	0.2	0.1	0.02	1.0	5.0	97.0	19.0	12.0	2.0	1.15	0.2	28
209449	2016-06-30 23:00:00	0.6	0.4	0.3	0.15	14.0	63.0	54.0	29.0	13.0	16.0	1.48	1.9	28
209454	2016-06-30 23:00:00	0.1	0.2	0.1	0.02	1.0	7.0	91.0	16.0	9.0	2.0	1.15	0.3	28
209473	2016-07-01 00:00:00	0.6	0.4	0.3	0.16	11.0	68.0	45.0	24.0	14.0	16.0	1.50	1.9	28
209478	2016-07-01 00:00:00	0.1	0.2	0.1	0.02	1.0	6.0	89.0	16.0	9.0	2.0	1.15	0.2	28

16932 rows × 14 columns



In [5]: `df.isnull().sum()`

```
Out[5]: date      0
      BEN      0
      CO      0
      EBE      0
      NMHC     0
      NO      0
      NO_2     0
      O_3      0
      PM10     0
      PM25     0
      SO_2     0
      TCH      0
      TOL      0
      station  0
      dtype: int64
```

In [6]: `df.describe()`

```
Out[6]:
```

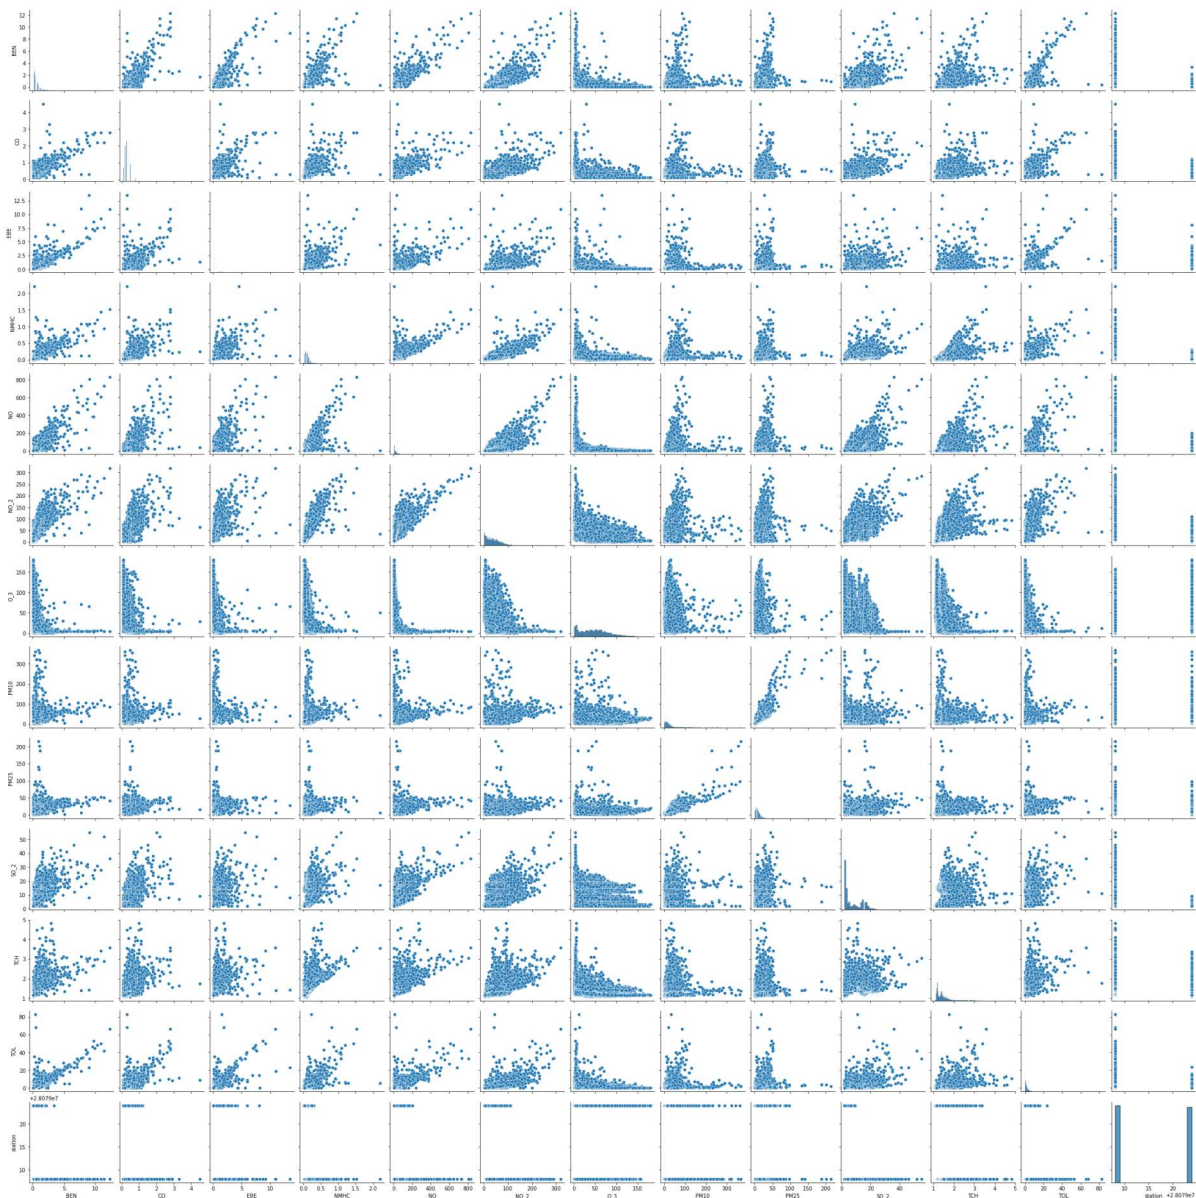
	BEN	CO	EBE	NMHC	NO	NO_2	
count	16932.000000	16932.000000	16932.000000	16932.000000	16932.000000	16932.000000	169
mean	0.537970	0.349941	0.298955	0.099913	20.815734	39.373376	
std	0.599479	0.203807	0.450204	0.079850	40.986063	31.170307	
min	0.100000	0.100000	0.100000	0.000000	1.000000	1.000000	
25%	0.200000	0.200000	0.100000	0.050000	1.000000	14.000000	
50%	0.400000	0.300000	0.200000	0.090000	7.000000	34.000000	
75%	0.700000	0.400000	0.300000	0.120000	23.000000	58.000000	
max	12.300000	4.500000	13.500000	2.210000	829.000000	319.000000	1

In [7]: `df.columns`

```
Out[7]: Index(['date', 'BEN', 'CO', 'EBE', 'NMHC', 'NO', 'NO_2', 'O_3', 'PM10', 'PM25',
              'SO_2', 'TCH', 'TOL', 'station'],
              dtype='object')
```

```
In [8]: sns.pairplot(df)
```

```
Out[8]: <seaborn.axisgrid.PairGrid at 0x20f4d2c7220>
```




```
In [9]: sns.distplot(df['PM10'])
```

C:\ProgramData\Anaconda3\lib\site-packages\seaborn\distributions.py:2557: FutureWarning: `distplot` is a deprecated function and will be removed in a future version. Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).


```
warnings.warn(msg, FutureWarning)
```

```
Out[9]: <AxesSubplot:xlabel='PM10', ylabel='Density'>
```



```
In [10]: sns.heatmap(df.corr())
```

```
Out[10]: <AxesSubplot:>
```



```
In [11]: df.loc[df['TCH']<2,'TCH']=0
df.loc[df['TCH']>2,'TCH']=1
df['TCH']=df['TCH'].astype(int)
df
```

C:\ProgramData\Anaconda3\lib\site-packages\pandas\core\indexing.py:1720: SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame.

Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy (https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)

```
self._setitem_single_column(loc, value, pi)
```

C:\ProgramData\Anaconda3\lib\site-packages\pandas\core\indexing.py:1720: SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame.

Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy (https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)

```
self._setitem_single_column(loc, value, pi)
```

<ipython-input-11-e3d36a273982>:3: SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame.

Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy (https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)

```
df['TCH']=df['TCH'].astype(int)
```


Out[11]:

	date	BEN	CO	EBE	NMHC	NO	NO_2	O_3	PM10	PM25	SO_2	TCH	TOL	
1	2016-11-01 01:00:00	3.1	1.1	2.0	0.53	260.0	144.0	4.0	46.0	24.0	18.0	1	14.4	28
6	2016-11-01 01:00:00	0.7	0.8	0.4	0.13	57.0	66.0	3.0	23.0	15.0	4.0	0	5.0	28
25	2016-11-01 02:00:00	2.7	1.0	2.1	0.40	139.0	114.0	4.0	37.0	21.0	14.0	1	15.0	28
30	2016-11-01 02:00:00	0.7	0.7	0.4	0.13	48.0	59.0	3.0	23.0	15.0	3.0	0	5.0	28
49	2016-11-01 03:00:00	1.7	0.8	1.4	0.25	53.0	90.0	4.0	31.0	19.0	10.0	0	10.7	28
...
209430	2016-06-30 22:00:00	0.1	0.2	0.1	0.02	1.0	5.0	97.0	19.0	12.0	2.0	0	0.2	28
209449	2016-06-30 23:00:00	0.6	0.4	0.3	0.15	14.0	63.0	54.0	29.0	13.0	16.0	0	1.9	28
209454	2016-06-30 23:00:00	0.1	0.2	0.1	0.02	1.0	7.0	91.0	16.0	9.0	2.0	0	0.3	28
209473	2016-07-01 00:00:00	0.6	0.4	0.3	0.16	11.0	68.0	45.0	24.0	14.0	16.0	0	1.9	28
209478	2016-07-01 00:00:00	0.1	0.2	0.1	0.02	1.0	6.0	89.0	16.0	9.0	2.0	0	0.2	28

16932 rows × 14 columns



LogisticRegression

```
In [12]: x=df[['BEN', 'CO', 'EBE', 'NMHC', 'NO', 'NO_2', 'O_3', 'PM10', 'PM25',
               'SO_2', 'TOL', 'station']]
y=df['TCH']
x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.3)
lgr=LogisticRegression()
lgr.fit(x_train,y_train)
```

Out[12]: LogisticRegression()

```
In [13]: lgr.predict(x_test)
```

```
Out[13]: array([0, 0, 0, ..., 0, 0, 0])
```

```
In [14]: lgr.score(x_test,y_test)
```

```
Out[14]: 0.9702755905511811
```

```
In [15]: fs=StandardScaler().fit_transform(x)
logr=LogisticRegression()
logr.fit(fs,y)
```

C:\ProgramData\Anaconda3\lib\site-packages\sklearn\linear_model_logistic.py:
763: ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max_iter) or scale the data as shown in:
<https://scikit-learn.org/stable/modules/preprocessing.html> (<https://scikit-learn.org/stable/modules/preprocessing.html>)
Please also refer to the documentation for alternative solver options:
https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression (https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression)
n_iter_i = _check_optimize_result(

```
Out[15]: LogisticRegression()
```

```
In [16]: o=[[1,2,3,4,5,6,7,8,9,10,11,12]]
prediction=logr.predict(o)
print(prediction)
```

```
[0]
```

```
In [17]: logr.classes_
```

```
Out[17]: array([0, 1, 2])
```

```
In [18]: logr.predict_proba(o)[0][0]
```

```
Out[18]: 0.99999999999999951
```

```
In [19]: logr.predict_proba(o)[0][1]
```

```
Out[19]: 4.962724005677087e-15
```

LinearRegression

```
In [20]: lr=LinearRegression()
lr.fit(x_train,y_train)
```

```
Out[20]: LinearRegression()
```

```
In [21]: print(lr.intercept_)
```

36395.73166804244


```
In [22]: coeff = pd.DataFrame(lr.coef_,x.columns,columns=['Co-efficient'])  
coeff
```

Out[22]:

	Co-efficient
BEN	-0.012937
CO	-0.099476
EBE	0.021893
NMHC	1.082968
NO	0.001672
NO_2	-0.001474
O_3	-0.000293
PM10	-0.000666
PM25	0.002112
SO_2	-0.003539
TOL	-0.000923
station	-0.001296

```
In [23]: prediction=lr.predict(x_test)  
plt.scatter(y_test,prediction)
```

Out[23]: <matplotlib.collections.PathCollection at 0x20f5df91370>



```
In [24]: print(lr.score(x_test,y_test))
```

0.31204558562684215

Ridge,Lasso

```
In [25]: rr=Ridge(alpha=10)
         rr.fit(x_train,y_train)
```

```
Out[25]: Ridge(alpha=10)
```

```
In [26]: rr.score(x_test,y_test)
```

```
Out[26]: 0.3113198549060787
```

```
In [27]: la=Lasso(alpha=10)
         la.fit(x_train,y_train)
```

```
Out[27]: Lasso(alpha=10)
```

```
In [28]: la.score(x_test,y_test)
```

```
Out[28]: -1.220663748746631e-06
```

ElasticNet

```
In [29]: en=ElasticNet()
         en.fit(x_train,y_train)
```

```
Out[29]: ElasticNet()
```

```
In [30]: print(en.coef_)
```

```
[ 0.          0.          0.          0.          0.00183732  0.
 -0.          0.          0.         -0.          0.         -0.         ]
```

```
In [31]: print(en.intercept_)
```

```
-0.007178338493906651
```

```
In [32]: print(en.predict(x_train))
```

```
[ 0.01119485 -0.00534102 -0.00534102 ...  0.00935753 -0.00534102
 -0.00534102]
```

```
In [33]: print(en.score(x_train,y_train))
```

```
0.24517361988217445
```

```
In [34]: print("Mean Absolytre Error:",metrics.mean_absolute_error(y_test,prediction))
```

```
Mean Absolytre Error: 0.057938190734859646
```

```
In [35]: print("Mean Square Error:", metrics.mean_squared_error(y_test, prediction))
```

Mean Square Error: 0.022627447265107833

```
In [36]: print("Root Mean Square Error:", np.sqrt(metrics.mean_absolute_error(y_test, prediction)))
```

Root Mean Square Error: 0.2407035328674252

RandomForest

```
In [37]: rfc=RandomForestClassifier()  
rfc.fit(x_train,y_train)
```

Out[37]: RandomForestClassifier()

```
In [38]: parameters={'max_depth':[1,2,3,4,5],  
                    'min_samples_leaf':[5,10,15,20,25],  
                    'n_estimators':[10,20,30,40,50]}
```

```
In [39]: grid_search=GridSearchCV(estimator=rfc,param_grid=parameters,cv=2,scoring="accuracy")  
grid_search.fit(x_train,y_train)
```

Out[39]: GridSearchCV(cv=2, estimator=RandomForestClassifier(),
param_grid={'max_depth': [1, 2, 3, 4, 5],
 'min_samples_leaf': [5, 10, 15, 20, 25],
 'n_estimators': [10, 20, 30, 40, 50]},
scoring='accuracy')

```
In [40]: grid_search.best_score_
```

Out[40]: 0.9785690178872763

```
In [41]: rfc_best=grid_search.best_estimator_
```

```
In [42]: plt.figure(figsize=(80,40))  
plot_tree(rfc_best.estimators_[5],class_names=['Yes','No','Yes','No'],filled=True)
```

```

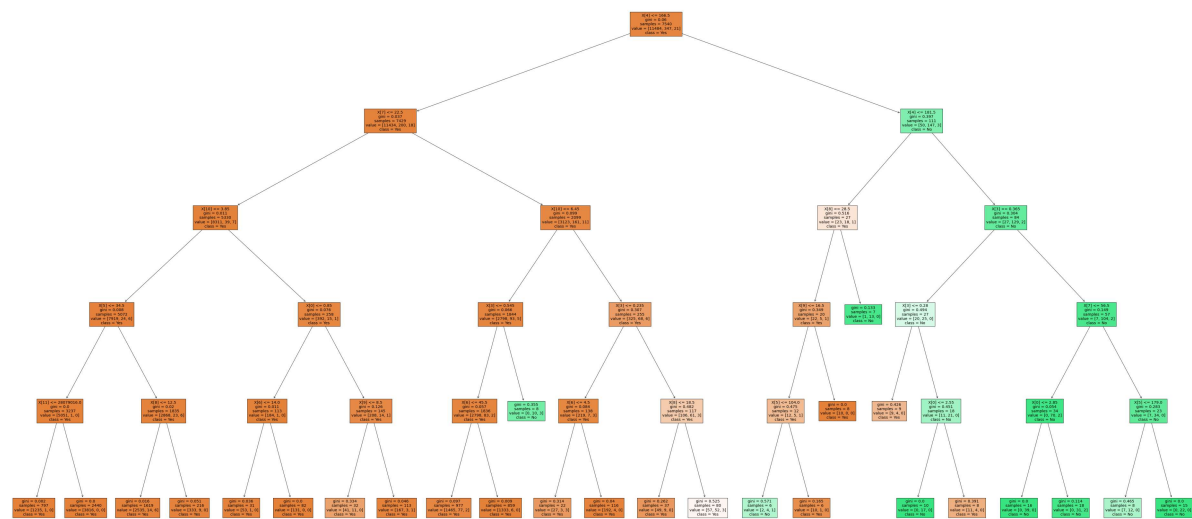
Out[42]: [Text(2426.086956521739, 1993.2, 'X[4] <= 166.5\ngini = 0.06\nsamples = 7540
\nvalue = [11484, 347, 21]\nnclass = Yes'),
Text(1431.391304347826, 1630.8000000000002, 'X[7] <= 22.5\ngini = 0.037\nsam
ples = 7429\nvalue = [11434, 200, 18]\nnclass = Yes'),
Text(776.3478260869565, 1268.4, 'X[10] <= 3.85\ngini = 0.011\nsamples = 5330
\nvalue = [8311, 39, 7]\nnclass = Yes'),
Text(388.17391304347825, 906.0, 'X[5] <= 34.5\ngini = 0.008\nsamples = 5072
\nvalue = [7919, 24, 6]\nnclass = Yes'),
Text(194.08695652173913, 543.5999999999999, 'X[11] <= 28079016.0\ngini = 0.0
\nsamples = 3237\nvalue = [5051, 1, 0]\nnclass = Yes'),
Text(97.04347826086956, 181.19999999999982, 'gini = 0.002\nsamples = 797\nva
lue = [1235, 1, 0]\nnclass = Yes'),
Text(291.1304347826087, 181.19999999999982, 'gini = 0.0\nsamples = 2440\nval
ue = [3816, 0, 0]\nnclass = Yes'),
Text(582.2608695652174, 543.5999999999999, 'X[8] <= 12.5\ngini = 0.02\nsampl
es = 1835\nvalue = [2868, 23, 6]\nnclass = Yes'),
Text(485.2173913043478, 181.19999999999982, 'gini = 0.016\nsamples = 1619\nv
alue = [2535, 14, 6]\nnclass = Yes'),
Text(679.304347826087, 181.19999999999982, 'gini = 0.051\nsamples = 216\nval
ue = [333, 9, 0]\nnclass = Yes'),
Text(1164.5217391304348, 906.0, 'X[0] <= 0.85\ngini = 0.076\nsamples = 258\n
value = [392, 15, 1]\nnclass = Yes'),
Text(970.4347826086956, 543.5999999999999, 'X[6] <= 14.0\ngini = 0.011\nsamp
les = 113\nvalue = [184, 1, 0]\nnclass = Yes'),
Text(873.391304347826, 181.19999999999982, 'gini = 0.036\nsamples = 31\nvalu
e = [53, 1, 0]\nnclass = Yes'),
Text(1067.4782608695652, 181.19999999999982, 'gini = 0.0\nsamples = 82\nvalu
e = [131, 0, 0]\nnclass = Yes'),
Text(1358.608695652174, 543.5999999999999, 'X[9] <= 8.5\ngini = 0.126\nsampl
es = 145\nvalue = [208, 14, 1]\nnclass = Yes'),
Text(1261.5652173913043, 181.19999999999982, 'gini = 0.334\nsamples = 32\nva
lue = [41, 11, 0]\nnclass = Yes'),
Text(1455.6521739130435, 181.19999999999982, 'gini = 0.046\nsamples = 113\nv
alue = [167, 3, 1]\nnclass = Yes'),
Text(2086.4347826086955, 1268.4, 'X[10] <= 6.45\ngini = 0.099\nsamples = 209
9\nvalue = [3123, 161, 11]\nnclass = Yes'),
Text(1843.8260869565217, 906.0, 'X[3] <= 0.545\ngini = 0.066\nsamples = 1844
\nvalue = [2798, 93, 5]\nnclass = Yes'),
Text(1746.782608695652, 543.5999999999999, 'X[6] <= 45.5\ngini = 0.057\nsamp
les = 1836\nvalue = [2798, 83, 2]\nnclass = Yes'),
Text(1649.7391304347825, 181.19999999999982, 'gini = 0.097\nsamples = 977\nv
alue = [1465, 77, 2]\nnclass = Yes'),
Text(1843.8260869565217, 181.19999999999982, 'gini = 0.009\nsamples = 859\nv
alue = [1333, 6, 0]\nnclass = Yes'),
Text(1940.8695652173913, 543.5999999999999, 'gini = 0.355\nsamples = 8\nvalu
e = [0, 10, 3]\nnclass = No'),
Text(2329.0434782608695, 906.0, 'X[3] <= 0.235\ngini = 0.307\nsamples = 255
\nvalue = [325, 68, 6]\nnclass = Yes'),
Text(2134.9565217391305, 543.5999999999999, 'X[6] <= 4.5\ngini = 0.084\nsamp
les = 138\nvalue = [219, 7, 3]\nnclass = Yes'),
Text(2037.9130434782608, 181.19999999999982, 'gini = 0.314\nsamples = 22\nva
lue = [27, 3, 3]\nnclass = Yes'),
Text(2232.0, 181.19999999999982, 'gini = 0.04\nsamples = 116\nvalue = [192,
4, 0]\nnclass = Yes'),
Text(2523.1304347826085, 543.5999999999999, 'X[8] <= 18.5\ngini = 0.482\nsam
ples = 117\nvalue = [106, 61, 3]\nnclass = Yes'),
Text(2426.086956521739, 181.19999999999982, 'gini = 0.262\nsamples = 37\nval

```

```

ue = [49, 9, 0]\n\nclass = Yes'),
Text(2620.173913043478, 181.19999999999982, 'gini = 0.525\n\ nsamples = 80\n\ nvalue = [57, 52, 3]\n\nclass = Yes'),
Text(3420.782608695652, 1630.8000000000002, 'X[4] <= 181.5\n\ ngini = 0.397\n\ nsamples = 111\n\ nvalue = [50, 147, 3]\n\nclass = No'),
Text(3105.391304347826, 1268.4, 'X[8] <= 28.5\n\ ngini = 0.516\n\ nsamples = 27\n\ nvalue = [23, 18, 1]\n\nclass = Yes'),
Text(3008.3478260869565, 906.0, 'X[9] <= 16.5\n\ ngini = 0.349\n\ nsamples = 20\n\ nvalue = [22, 5, 1]\n\nclass = Yes'),
Text(2911.304347826087, 543.5999999999999, 'X[5] <= 104.0\n\ ngini = 0.475\n\ nsamples = 12\n\ nvalue = [12, 5, 1]\n\nclass = Yes'),
Text(2814.2608695652175, 181.19999999999982, 'gini = 0.571\n\ nsamples = 6\n\ nvalue = [2, 4, 1]\n\nclass = No'),
Text(3008.3478260869565, 181.19999999999982, 'gini = 0.165\n\ nsamples = 6\n\ nvalue = [10, 1, 0]\n\nclass = Yes'),
Text(3105.391304347826, 543.5999999999999, 'gini = 0.0\n\ nsamples = 8\n\ nvalue = [10, 0, 0]\n\nclass = Yes'),
Text(3202.4347826086955, 906.0, 'gini = 0.133\n\ nsamples = 7\n\ nvalue = [1, 13, 0]\n\nclass = No'),
Text(3736.173913043478, 1268.4, 'X[3] <= 0.365\n\ ngini = 0.304\n\ nsamples = 84\n\ nvalue = [27, 129, 2]\n\nclass = No'),
Text(3396.5217391304345, 906.0, 'X[3] <= 0.28\n\ ngini = 0.494\n\ nsamples = 27\n\ nvalue = [20, 25, 0]\n\nclass = No'),
Text(3299.478260869565, 543.5999999999999, 'gini = 0.426\n\ nsamples = 9\n\ nvalue = [9, 4, 0]\n\nclass = Yes'),
Text(3493.565217391304, 543.5999999999999, 'X[0] <= 2.55\n\ ngini = 0.451\n\ nsamples = 18\n\ nvalue = [11, 21, 0]\n\nclass = No'),
Text(3396.5217391304345, 181.19999999999982, 'gini = 0.0\n\ nsamples = 10\n\ nvalue = [0, 17, 0]\n\nclass = No'),
Text(3590.608695652174, 181.19999999999982, 'gini = 0.391\n\ nsamples = 8\n\ nvalue = [11, 4, 0]\n\nclass = Yes'),
Text(4075.8260869565215, 906.0, 'X[7] <= 56.5\n\ ngini = 0.149\n\ nsamples = 57\n\ nvalue = [7, 104, 2]\n\nclass = No'),
Text(3881.7391304347825, 543.5999999999999, 'X[0] <= 2.85\n\ ngini = 0.054\n\ nsamples = 34\n\ nvalue = [0, 70, 2]\n\nclass = No'),
Text(3784.695652173913, 181.19999999999982, 'gini = 0.0\n\ nsamples = 16\n\ nvalue = [0, 39, 0]\n\nclass = No'),
Text(3978.782608695652, 181.19999999999982, 'gini = 0.114\n\ nsamples = 18\n\ nvalue = [0, 31, 2]\n\nclass = No'),
Text(4269.913043478261, 543.5999999999999, 'X[5] <= 179.0\n\ ngini = 0.283\n\ nsamples = 23\n\ nvalue = [7, 34, 0]\n\nclass = No'),
Text(4172.869565217391, 181.19999999999982, 'gini = 0.465\n\ nsamples = 8\n\ nvalue = [7, 12, 0]\n\nclass = No'),
Text(4366.95652173913, 181.19999999999982, 'gini = 0.0\n\ nsamples = 15\n\ nvalue = [0, 22, 0]\n\nclass = No')]

```

Best model:LogisticRegression

In []: