

```
In [1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.linear_model import Ridge,Lasso
from sklearn.linear_model import ElasticNet
from sklearn import metrics
from sklearn.linear_model import LogisticRegression
from sklearn.preprocessing import StandardScaler
from sklearn.ensemble import RandomForestClassifier
from sklearn.model_selection import GridSearchCV
from sklearn.tree import plot_tree
```

```
In [2]: df=pd.read_csv(r"C:\Users\user\Downloads\csvs_per_year\csvs_per_year\madrid_2018\
df
```

Out[2]:

	date	BEN	CH4	CO	EBE	NMHC	NO	NO_2	NOx	O_3	PM10	PM25	SO_2	T
0	2018-03-01 01:00:00	NaN	NaN	0.3	NaN	NaN	1.0	29.0	31.0	NaN	NaN	NaN	2.0	NaN
1	2018-03-01 01:00:00	0.5	1.39	0.3	0.2	0.02	6.0	40.0	49.0	52.0	5.0	4.0	3.0	1
2	2018-03-01 01:00:00	0.4	NaN	NaN	0.2	NaN	4.0	41.0	47.0	NaN	NaN	NaN	NaN	NaN
3	2018-03-01 01:00:00	NaN	NaN	0.3	NaN	NaN	1.0	35.0	37.0	54.0	NaN	NaN	NaN	NaN
4	2018-03-01 01:00:00	NaN	NaN	NaN	NaN	NaN	1.0	27.0	29.0	49.0	NaN	NaN	3.0	NaN
...	...	...	...	...	...	...	...	...	...	...	...	...	...	...
69091	2018-02-01 00:00:00	NaN	NaN	0.5	NaN	NaN	66.0	91.0	192.0	1.0	35.0	22.0	NaN	NaN
69092	2018-02-01 00:00:00	NaN	NaN	0.7	NaN	NaN	87.0	107.0	241.0	NaN	29.0	NaN	15.0	NaN
69093	2018-02-01 00:00:00	NaN	NaN	NaN	NaN	NaN	28.0	48.0	91.0	2.0	NaN	NaN	NaN	NaN
69094	2018-02-01 00:00:00	NaN	NaN	NaN	NaN	NaN	141.0	103.0	320.0	2.0	NaN	NaN	NaN	NaN
69095	2018-02-01 00:00:00	NaN	NaN	NaN	NaN	NaN	69.0	96.0	202.0	3.0	26.0	NaN	NaN	NaN

69096 rows × 16 columns



```
In [3]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 69096 entries, 0 to 69095
Data columns (total 16 columns):
#   Column      Non-Null Count  Dtype
---  -
0   date        69096 non-null  object
1   BEN         16950 non-null  float64
2   CH4         8440 non-null   float64
3   CO          28598 non-null  float64
4   EBE         16949 non-null  float64
5   NMHC        8440 non-null   float64
6   NO          68826 non-null  float64
7   NO_2        68826 non-null  float64
8   NOx         68826 non-null  float64
9   O_3         40049 non-null  float64
10  PM10        36911 non-null  float64
11  PM25        18912 non-null  float64
12  SO_2        28586 non-null  float64
13  TCH         8440 non-null   float64
14  TOL         16950 non-null  float64
15  station     69096 non-null  int64
dtypes: float64(14), int64(1), object(1)
memory usage: 8.4+ MB
```

```
In [4]: df=df.dropna()  
df
```

Out[4]:

	date	BEN	CH4	CO	EBE	NMHC	NO	NO_2	NOx	O_3	PM10	PM25	SO_2	TC
1	2018-03-01 01:00:00	0.5	1.39	0.3	0.2	0.02	6.0	40.0	49.0	52.0	5.0	4.0	3.0	1.0
6	2018-03-01 01:00:00	0.4	1.11	0.2	0.1	0.06	1.0	25.0	27.0	55.0	5.0	4.0	4.0	1.0
25	2018-03-01 02:00:00	0.4	1.42	0.2	0.1	0.01	4.0	26.0	32.0	64.0	4.0	4.0	3.0	1.0
30	2018-03-01 02:00:00	0.3	1.10	0.2	0.1	0.05	1.0	12.0	13.0	69.0	5.0	4.0	4.0	1.0
49	2018-03-01 03:00:00	0.3	1.41	0.2	0.1	0.01	3.0	16.0	20.0	68.0	3.0	2.0	3.0	1.0
...	...	...	...	...	...	...	...	...	...	...	...	...	...	...
69030	2018-01-31 22:00:00	1.8	1.21	0.7	1.7	0.19	151.0	129.0	361.0	1.0	45.0	26.0	11.0	1.0
69049	2018-01-31 23:00:00	3.1	1.87	1.2	2.0	0.35	296.0	162.0	615.0	3.0	39.0	23.0	8.0	2.0
69054	2018-01-31 23:00:00	1.6	1.17	0.6	1.4	0.15	127.0	106.0	301.0	1.0	43.0	25.0	8.0	1.0
69073	2018-02-01 00:00:00	3.2	1.53	1.0	2.1	0.19	125.0	117.0	309.0	3.0	37.0	24.0	6.0	1.0
69078	2018-02-01 00:00:00	1.3	1.14	0.4	0.8	0.10	54.0	73.0	155.0	1.0	27.0	16.0	5.0	1.0

4562 rows × 16 columns

In [5]: `df.isnull().sum()`

```
Out[5]: date      0
        BEN      0
        CH4      0
        CO       0
        EBE      0
        NMHC     0
        NO       0
        NO_2     0
        NOx      0
        O_3      0
        PM10     0
        PM25     0
        SO_2     0
        TCH      0
        TOL      0
        station  0
        dtype: int64
```

In [6]: `df.describe()`

```
Out[6]:
```

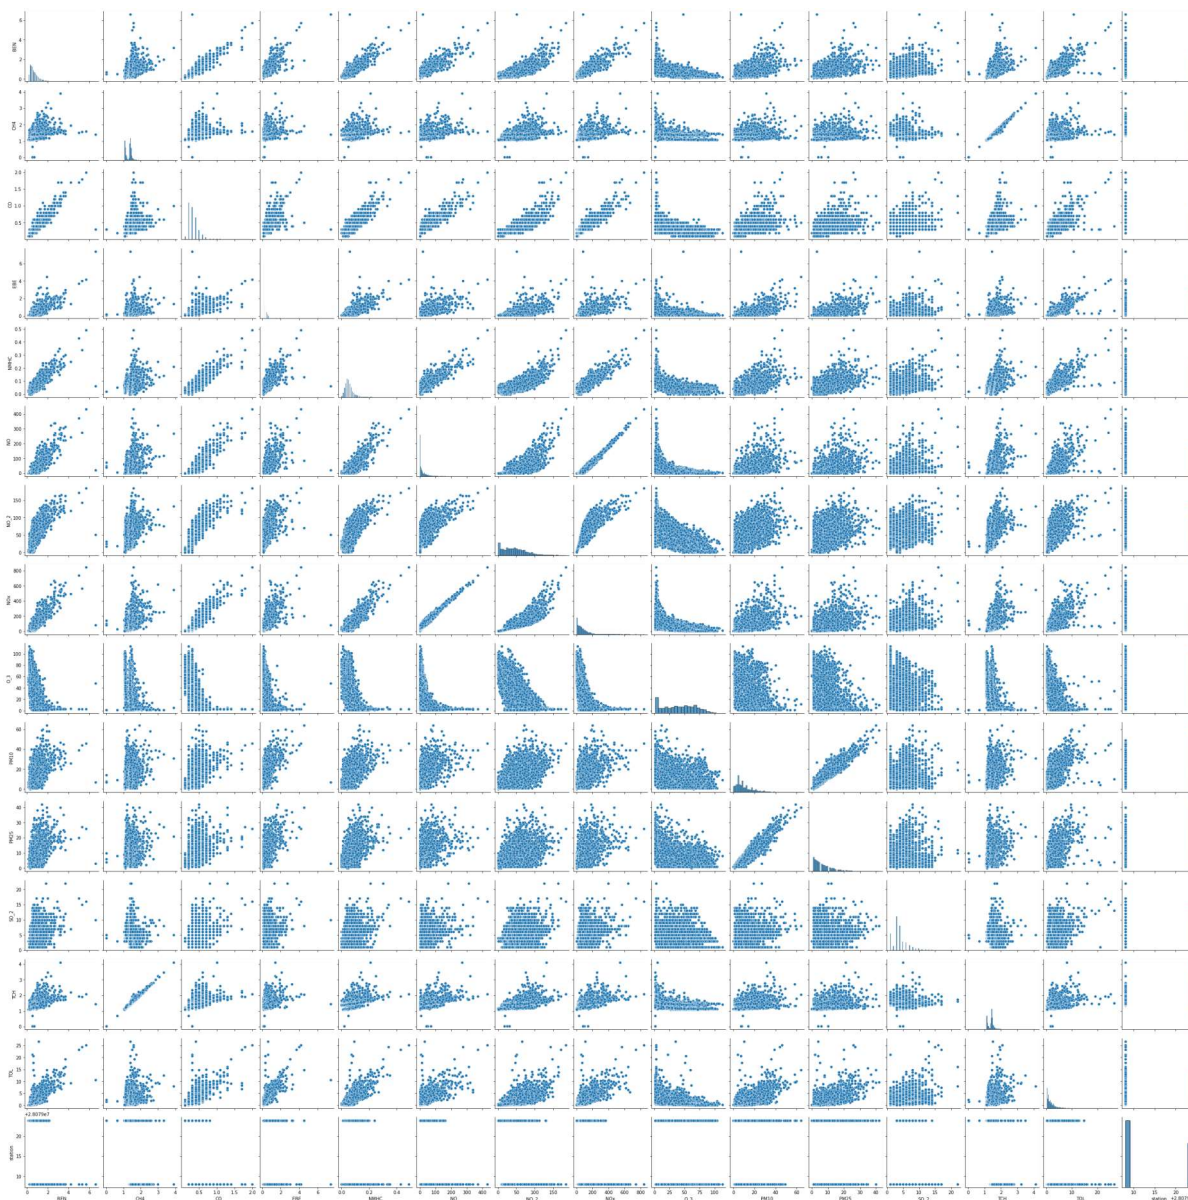
	BEN	CH4	CO	EBE	NMHC	NO	NO_2
<b>count</b>	4562.00000	4562.000000	4562.000000	4562.000000	4562.000000	4562.000000	4562.000000
<b>mean</b>	0.69349	1.329163	0.330579	0.286782	0.056773	21.742218	44.152126
<b>std</b>	0.46832	0.214399	0.161489	0.354442	0.037711	35.539531	30.234015
<b>min</b>	0.10000	0.020000	0.100000	0.100000	0.000000	1.000000	1.000000
<b>25%</b>	0.40000	1.120000	0.200000	0.100000	0.030000	1.000000	20.000000
<b>50%</b>	0.60000	1.390000	0.300000	0.200000	0.050000	9.000000	41.000000
<b>75%</b>	0.90000	1.420000	0.400000	0.300000	0.070000	27.000000	64.000000
<b>max</b>	6.60000	3.920000	2.000000	7.400000	0.490000	431.000000	184.000000

In [7]: `df.columns`

```
Out[7]: Index(['date', 'BEN', 'CH4', 'CO', 'EBE', 'NMHC', 'NO', 'NO_2', 'NOx', 'O_3',
              'PM10', 'PM25', 'SO_2', 'TCH', 'TOL', 'station'],
              dtype='object')
```

```
In [8]: sns.pairplot(df)
```

```
Out[8]: <seaborn.axisgrid.PairGrid at 0x21ddf7180a0>
```




```
In [9]: sns.distplot(df['PM10'])
```

C:\ProgramData\Anaconda3\lib\site-packages\seaborn\distributions.py:2557: FutureWarning: `distplot` is a deprecated function and will be removed in a future version. Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).


```
warnings.warn(msg, FutureWarning)
```

```
Out[9]: <AxesSubplot:xlabel='PM10', ylabel='Density'>
```



```
In [10]: sns.heatmap(df.corr())
```

```
Out[10]: <AxesSubplot:>
```



```
In [11]: df.loc[df['TCH']<2,'TCH']=0
df.loc[df['TCH']>2,'TCH']=1
df['TCH']=df['TCH'].astype(int)
df
```

C:\ProgramData\Anaconda3\lib\site-packages\pandas\core\indexing.py:1720: SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame.

Try using .loc[row\_indexer,col\_indexer] = value instead

See the caveats in the documentation: [https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy) ([https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy))

```
self._setitem_single_column(loc, value, pi)
```

C:\ProgramData\Anaconda3\lib\site-packages\pandas\core\indexing.py:1720: SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame.

Try using .loc[row\_indexer,col\_indexer] = value instead

See the caveats in the documentation: [https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy) ([https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy))

```
self._setitem_single_column(loc, value, pi)
```

<ipython-input-11-e3d36a273982>:3: SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame.

Try using .loc[row\_indexer,col\_indexer] = value instead

See the caveats in the documentation: [https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy) ([https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy))

```
df['TCH']=df['TCH'].astype(int)
```



Out[11]:

	date	BEN	CH4	CO	EBE	NMHC	NO	NO_2	NOx	O_3	PM10	PM25	SO_2	TC
1	2018-03-01 01:00:00	0.5	1.39	0.3	0.2	0.02	6.0	40.0	49.0	52.0	5.0	4.0	3.0	
6	2018-03-01 01:00:00	0.4	1.11	0.2	0.1	0.06	1.0	25.0	27.0	55.0	5.0	4.0	4.0	
25	2018-03-01 02:00:00	0.4	1.42	0.2	0.1	0.01	4.0	26.0	32.0	64.0	4.0	4.0	3.0	
30	2018-03-01 02:00:00	0.3	1.10	0.2	0.1	0.05	1.0	12.0	13.0	69.0	5.0	4.0	4.0	
49	2018-03-01 03:00:00	0.3	1.41	0.2	0.1	0.01	3.0	16.0	20.0	68.0	3.0	2.0	3.0	
...	...	...	...	...	...	...	...	...	...	...	...	...	...	...
69030	2018-01-31 22:00:00	1.8	1.21	0.7	1.7	0.19	151.0	129.0	361.0	1.0	45.0	26.0	11.0	
69049	2018-01-31 23:00:00	3.1	1.87	1.2	2.0	0.35	296.0	162.0	615.0	3.0	39.0	23.0	8.0	
69054	2018-01-31 23:00:00	1.6	1.17	0.6	1.4	0.15	127.0	106.0	301.0	1.0	43.0	25.0	8.0	
69073	2018-02-01 00:00:00	3.2	1.53	1.0	2.1	0.19	125.0	117.0	309.0	3.0	37.0	24.0	6.0	
69078	2018-02-01 00:00:00	1.3	1.14	0.4	0.8	0.10	54.0	73.0	155.0	1.0	27.0	16.0	5.0	

4562 rows × 16 columns



## LogisticRegression

```
In [12]: x=df[['BEN', 'CO', 'EBE', 'NMHC', 'NO', 'NO_2', 'O_3', 'PM10', 'PM25',
               'SO_2', 'TOL', 'station']]
y=df['TCH']
x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.3)
lgr=LogisticRegression()
lgr.fit(x_train,y_train)
```

Out[12]: LogisticRegression()

```
In [13]: lgr.predict(x_test)
```

```
Out[13]: array([0, 0, 0, ..., 0, 0, 0])
```

```
In [14]: lgr.score(x_test,y_test)
```

```
Out[14]: 0.9824689554419284
```

```
In [15]: fs=StandardScaler().fit_transform(x)
logr=LogisticRegression()
logr.fit(fs,y)
```

C:\ProgramData\Anaconda3\lib\site-packages\sklearn\linear\_model\\_logistic.py:  
763: ConvergenceWarning: lbfgs failed to converge (status=1):  
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max\_iter) or scale the data as shown in:  
<https://scikit-learn.org/stable/modules/preprocessing.html> (<https://scikit-learn.org/stable/modules/preprocessing.html>)  
Please also refer to the documentation for alternative solver options:  
[https://scikit-learn.org/stable/modules/linear\\_model.html#logistic-regression](https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression) ([https://scikit-learn.org/stable/modules/linear\\_model.html#logistic-regression](https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression))  
n\_iter\_i = \_check\_optimize\_result(

```
Out[15]: LogisticRegression()
```

```
In [16]: o=[[1,2,3,4,5,6,7,8,9,10,11,12]]
prediction=logr.predict(o)
print(prediction)
```

```
[0]
```

```
In [17]: logr.classes_
```

```
Out[17]: array([0, 1, 2])
```

```
In [18]: logr.predict_proba(o)[0][0]
```

```
Out[18]: 0.9999999975037841
```

```
In [19]: logr.predict_proba(o)[0][1]
```

```
Out[19]: 5.085510314304287e-16
```

## LinearRegression

```
In [20]: lr=LinearRegression()
lr.fit(x_train,y_train)
```

```
Out[20]: LinearRegression()
```

In [21]: `print(lr.intercept_)`

104579.50711627983


In [22]: `coeff = pd.DataFrame(lr.coef_,x.columns,columns=['Co-efficient'])`  
`coeff`

Out[22]:

	Co-efficient
<b>BEN</b>	-0.016616
<b>CO</b>	-0.396203
<b>EBE</b>	0.061732
<b>NMHC</b>	1.052382
<b>NO</b>	0.002328
<b>NO_2</b>	-0.000613
<b>O_3</b>	-0.000199
<b>PM10</b>	-0.001645
<b>PM25</b>	0.003811
<b>SO_2</b>	-0.002612
<b>TOL</b>	-0.005893
<b>station</b>	-0.003724

In [23]: `prediction=lr.predict(x_test)`  
`plt.scatter(y_test,prediction)`

Out[23]: <matplotlib.collections.PathCollection at 0x21df29904f0>



In [24]: `print(lr.score(x_test,y_test))`

0.16185028565714

## Ridge,Lasso

```
In [25]: rr=Ridge(alpha=10)
rr.fit(x_train,y_train)
```

```
Out[25]: Ridge(alpha=10)
```

```
In [26]: rr.score(x_test,y_test)
```

```
Out[26]: 0.15748597817418575
```

```
In [27]: la=Lasso(alpha=10)
la.fit(x_train,y_train)
```

```
Out[27]: Lasso(alpha=10)
```

```
In [28]: la.score(x_test,y_test)
```

```
Out[28]: -0.00020661188397363972
```

## ElasticNet

```
In [29]: en=ElasticNet()
en.fit(x_train,y_train)
```

```
Out[29]: ElasticNet()
```

```
In [30]: print(en.coef_)
```

```
[ 0.         0.         0.         0.         0.00122064  0.
 -0.         0.         0.        -0.         0.         0.        ]
```

```
In [31]: print(en.intercept_)
```

```
-0.007178926900212195
```

```
In [32]: print(en.predict(x_train))
```

```
[-0.00595829 -0.00595829  0.0807071  ...  0.02211641  0.00258619
  0.05873559]
```

```
In [33]: print(en.score(x_train,y_train))
```

```
0.14800264044790312
```

```
In [34]: print("Mean Absolytre Error:",metrics.mean_absolute_error(y_test,prediction))
```

```
Mean Absolytre Error: 0.04467923237090176
```

```
In [35]: print("Mean Square Error:", metrics.mean_squared_error(y_test, prediction))
```

Mean Square Error: 0.014436045131121353

```
In [36]: print("Root Mean Square Error:", np.sqrt(metrics.mean_absolute_error(y_test, prediction)))
```

Root Mean Square Error: 0.21137462565526108

## RandomForest

```
In [37]: rfc=RandomForestClassifier()  
rfc.fit(x_train,y_train)
```

Out[37]: RandomForestClassifier()

```
In [38]: parameters={'max_depth':[1,2,3,4,5],  
                    'min_samples_leaf':[5,10,15,20,25],  
                    'n_estimators':[10,20,30,40,50]}
```

```
In [39]: grid_search=GridSearchCV(estimator=rfc,param_grid=parameters,cv=2,scoring="accuracy")  
grid_search.fit(x_train,y_train)
```

Out[39]: GridSearchCV(cv=2, estimator=RandomForestClassifier(),  
param\_grid={'max\_depth': [1, 2, 3, 4, 5],  
            'min\_samples\_leaf': [5, 10, 15, 20, 25],  
            'n\_estimators': [10, 20, 30, 40, 50]},  
scoring='accuracy')

```
In [40]: grid_search.best_score_
```

Out[40]: 0.9821483498979132

```
In [41]: rfc_best=grid_search.best_estimator_
```

```
In [42]: plt.figure(figsize=(80,40))  
plot_tree(rfc_best.estimators_[5],class_names=['Yes','No','Yes','No'],filled=True)
```

```

Out[42]: [Text(2466.9473684210525, 1993.2, 'X[5] <= 78.5\ngini = 0.037\nsamples = 1990\nvalue = [3133, 56, 4]\nclass = Yes'),
Text(1409.6842105263158, 1630.8000000000002, 'X[4] <= 99.0\ngini = 0.005\nsamples = 1723\nvalue = [2764, 6, 1]\nclass = Yes'),
Text(1174.7368421052631, 1268.4, 'X[8] <= 20.5\ngini = 0.004\nsamples = 1718\nvalue = [2760, 4, 1]\nclass = Yes'),
Text(704.8421052631579, 906.0, 'X[4] <= 85.5\ngini = 0.001\nsamples = 1676\nvalue = [2686, 1, 1]\nclass = Yes'),
Text(469.89473684210526, 543.5999999999999, 'X[5] <= 69.5\ngini = 0.001\nsamples = 1671\nvalue = [2680, 1, 0]\nclass = Yes'),
Text(234.94736842105263, 181.19999999999982, 'gini = 0.0\nsamples = 1547\nvalue = [2479, 0, 0]\nclass = Yes'),
Text(704.8421052631579, 181.19999999999982, 'gini = 0.01\nsamples = 124\nvalue = [201, 1, 0]\nclass = Yes'),
Text(939.7894736842105, 543.5999999999999, 'gini = 0.245\nsamples = 5\nvalue = [6, 0, 1]\nclass = Yes'),
Text(1644.6315789473683, 906.0, 'X[1] <= 0.45\ngini = 0.075\nsamples = 42\nvalue = [74, 3, 0]\nclass = Yes'),
Text(1409.6842105263158, 543.5999999999999, 'gini = 0.0\nsamples = 30\nvalue = [57, 0, 0]\nclass = Yes'),
Text(1879.578947368421, 543.5999999999999, 'X[6] <= 6.5\ngini = 0.255\nsamples = 12\nvalue = [17, 3, 0]\nclass = Yes'),
Text(1644.6315789473683, 181.19999999999982, 'gini = 0.408\nsamples = 5\nvalue = [5, 2, 0]\nclass = Yes'),
Text(2114.5263157894738, 181.19999999999982, 'gini = 0.142\nsamples = 7\nvalue = [12, 1, 0]\nclass = Yes'),
Text(1644.6315789473683, 1268.4, 'gini = 0.444\nsamples = 5\nvalue = [4, 2, 0]\nclass = Yes'),
Text(3524.2105263157896, 1630.8000000000002, 'X[2] <= 2.05\ngini = 0.221\nsamples = 267\nvalue = [369, 50, 3]\nclass = Yes'),
Text(3054.315789473684, 1268.4, 'X[1] <= 0.45\ngini = 0.189\nsamples = 257\nvalue = [363, 43, 0]\nclass = Yes'),
Text(2819.3684210526317, 906.0, 'gini = 0.0\nsamples = 51\nvalue = [86, 0, 0]\nclass = Yes'),
Text(3289.2631578947367, 906.0, 'X[4] <= 159.5\ngini = 0.233\nsamples = 206\nvalue = [277, 43, 0]\nclass = Yes'),
Text(2819.3684210526317, 543.5999999999999, 'X[0] <= 1.15\ngini = 0.173\nsamples = 185\nvalue = [255, 27, 0]\nclass = Yes'),
Text(2584.4210526315787, 181.19999999999982, 'gini = 0.047\nsamples = 51\nvalue = [81, 2, 0]\nclass = Yes'),
Text(3054.315789473684, 181.19999999999982, 'gini = 0.22\nsamples = 134\nvalue = [174, 25, 0]\nclass = Yes'),
Text(3759.157894736842, 543.5999999999999, 'X[4] <= 199.0\ngini = 0.488\nsamples = 21\nvalue = [22, 16, 0]\nclass = Yes'),
Text(3524.2105263157896, 181.19999999999982, 'gini = 0.444\nsamples = 12\nvalue = [7, 14, 0]\nclass = No'),
Text(3994.1052631578946, 181.19999999999982, 'gini = 0.208\nsamples = 9\nvalue = [15, 2, 0]\nclass = Yes'),
Text(3994.1052631578946, 1268.4, 'X[0] <= 2.95\ngini = 0.633\nsamples = 10\nvalue = [6, 7, 3]\nclass = No'),
Text(3759.157894736842, 906.0, 'gini = 0.54\nsamples = 5\nvalue = [1, 6, 3]\nclass = No'),
Text(4229.0526315789475, 906.0, 'gini = 0.278\nsamples = 5\nvalue = [5, 1, 0]\nclass = Yes')]

```

